

# Lecture 14: Deadlocks

Professor G. Sandoval

Some slides derived from : Tanenbaum/Bo, John Regehr, and Brendan Dolan-Gavitt  
Thanks !!

# Today

→ • Xv6 Locks

- Deadlocks

# XV6 Locks

- XV6 represents a lock a `struct spinlock`.
- The critical field in the struct is `locked` {0 = available, non-zero = being held}

## Spinlock.h

```
00001: // Mutual exclusion lock.
00002: struct spinlock {
00003:     uint locked;           // Is the lock held?
00004:
00005:     // For debugging:
```

# XV6 Locks

- To acquire a lock you call `acquire`, which is implemented as follows:

```
21  void
22  acquire(struct spinlock *lk)
23  {
24      for(;;) {
25          if(!lk->locked) {
26              lk->locked = 1;
27              break;
28          }
29      }
30  }
```

# XV6 Locks

- To acquire a lock you call `acquire`, which is implemented as follows:

```
21  void
22  acquire(struct spinlock *lk)
23  {
24      for(;;) {
25          if(!lk->locked) {
26              lk->locked = 1;
27              break;
28          }
29      }
30  }
```

# XV6 Locks

- To acquire a lock you call **acquire**, which is really implemented as follows:

```
00020: // Acquire the lock.
00021: // Loops (spins) until the lock is acquired.
00022: // Holding a lock for a long time may cause
00023: // other CPUs to waste time spinning to acquire it.
00024: void
00025: acquire(struct spinlock *lk)
00026: {
00027:     pushcli(); // disable interrupts to avoid deadlock.
00028:     if(holding(lk))
00029:         panic("acquire");
00030:
00031:     // The xchg is atomic.
00032:     // It also serializes, so that reads after acquire are not
00033:     // reordered before it.
00034:     while(xchg(&lk->locked, 1) != 0)
00035:         ;
00036:
00037:     // Record info about lock acquisition for debugging.
00038:     lk->cpu = cpu;
00039:     getcallerpcs(&lk, lk->pcs);
00040: }
```

# XV6 Locks

- It's counterpart **release** is implemented as follows:

```
00042: // Release the lock.
00043: void
00044: release(struct spinlock *lk)
00045: {
00046:     if(!holding(lk))
00047:         panic("release");
00048:
00049:     lk->pcs[0] = 0;
00050:     lk->cpu = 0;
00051:
00052:     xchg(&lk->locked, 0);
00053:
00054:     popcli();
00055: }
```

# XV6 Locks Usage

```
00414: // Kill the process with the given pid.
00415: // Process won't exit until it returns
00416: // to user space (see trap in trap.c).
00417: int
00418: kill(int pid)
00419: {
00420:     struct proc *p;
00421:
00422:     acquire(&ptable.lock);
00423:     for(p = ptable.proc; p < &ptable.proc[NPROC]; p++){
00424:         if(p->pid == pid){
00425:             p->killed = 1;
00426:             // Wake process from sleep if necessary.
00427:             if(p->state == SLEEPING)
00428:                 p->state = RUNNABLE;
00429:             release(&ptable.lock);
00430:             return 0;
00431:         }
00432:     }
00433:     release(&ptable.lock);
00434:     return -1;
00435: }
```



# Today

- XV6 Locks

- • Deadlocks

# Dining Philosophers

```
#define N 5                                     /* number of philosophers */

void philosopher(int i)                         /* i: philosopher number, from 0 to 4 */
{
    while (TRUE) {
        think();                               /* philosopher is thinking */
        take_fork(i);                          /* take left fork */
        take_fork((i+1) % N);                  /* take right fork; % is modulo operator */
        eat();                                 /* yum-yum, spaghetti */
        put_fork(i);                          /* put left fork back on the table */
        put_fork((i+1) % N);                  /* put right fork back on the table */
    }
}
```

Solution to the dining philosophers problem?

# Deadlock

- Deadlock: two or more threads are waiting on events that only those threads can generate.
  - Computer don't see the big picture needed to break the stalemate.
  - Hard to handle automatically -> lots of theory, little practice.
- Livelock: thread blocked indefinitely by other thread(s) using a resource.
- Livelock naturally goes away when system load decreases
- Dealock does not go away by itself





# DEADLOCK

Game over, man, game over.

# Required Conditions for Deadlock

- Mutual Exclusion: Resources cannot be shared.
- Hold and Wait: A thread is both holding a resource and waiting on another resource to become free.
- No preemption: once a thread gets a resource, it cannot be taken away.
- Circular Wait: There is a cycle in the graph of who has and who wants what.



- What's the smallest number of threads needed to deadlock?



# Dealing with Deadlock

- **Deadlock ignorance**

- Why worry?

- **Deadlock detection**

- Figure out when deadlock has occurred and “deal with it”
  - » Figuring it out → mildly difficult
  - » Dealing with it → often messy, may need to reboot

- **Deadlock avoidance**

- Reject resource requests that might lead to deadlock

- **Deadlock prevention**

- Use “rules” that make it impossible for all four conditions to hold



# Deadlock Detection

- **Resource allocation graph**

- **Resources**  $\{r_1, \dots, r_m\}$
- **Threads**  $\{t_1, \dots, t_m\}$
- **Edges:**
  - »  $t_i \rightarrow r_j$ :  $t_i$  wants  $r_j$
  - »  $r_j \rightarrow t_i$ :  $r_j$  allocated to  $t_i$


- **Acyclic  $\rightarrow$  no deadlocks**

- **Cyclic  $\rightarrow$  deadlock**

- **Cycle detection**

- **Several known algorithms**
- **$O(n^2)$ ,  $n = |T| + |R|$**

- **Alternative approach: Wait for someone to hold a resource for way to long, then decide that deadlock has occurred**





# Deadlock Detection


- Periodically scan for deadlocks

- When to scan:

- Just before granting resources?
- Whenever resources unavailable?
- Fixed intervals?
- When CPU utilization drops?
- When thread not runnable for extended period of time?

- How to recover?


- Terminate threads
- Break locks
- Invoke exception handlers
- Create more resources
- Reboot
- ...



# Deadlock Avoidance

## ● Idea:

- Run version of deadlock detection algorithm in resource allocator
- Add “claim” edges to resources threads may request
- A cycle in extended graph  $\rightarrow$  unsafe state  $\rightarrow$  may lead to deadlock
- Deny allocations that result in unsafe states
  - » Claim edge converted to request edge
  - » Thread blocks until request can be safely satisfied



Do you think this is used often?  
Why or why not?

# When to worry?

- When do you have to worry about deadlock?
- Any time you write code that acquires more than one lock

# Approaches in Practice

- User code in Linux
  - Helgrind – a Valgrind plugin that checks for circular wait.
- Linux Kernel
  - Lock validator – dynamically checks for circular wait
- Windows thread pool
  - Create more threads if stuck
- Linux CPU Scheduler
  - Avoid starving any thread.

# Important from Today

- **Four conditions required for deadlock to occur:**
  - Mutual exclusion
  - Hold and wait
  - No resource pre-emption
  - Circular wait
- **Four techniques for handling deadlock**
  - Prevention: Design rules so one condition cannot occur
  - Avoidance: Dynamically deny “unsafe” requests
  - Detection and recovery: Let it happen, notice it, and fix it
  - Ignore: Do nothing, hope it does not happen, reboot often