

Lecture 4: Assembly

Professor G. Sandoval

Some slides adapted by G. Sandoval for CS3224, from Tanenbaum & Bo, Modern Operating Systems: 4th ed., (c) 2013 Prentice-Hall, Inc. All rights reserved.
Also some Slides by Brendan Dolan-Gavitt and Bryant and O'Hallaron, Computer Systems: A programmer's Perspective, Third Edition

Today: Machine Programming Basics

- History of Intel processors and architectures
- C, assembly, machine code
- Assembly Basics: Registers, operands, move
- Arithmetic & logical operations

The PC



Intel x86 Processors


- **Dominate laptop/desktop/server market**
- **Evolutionary design**
 - Backwards compatible up until 8086, introduced in 1978
 - Added more features as time goes on
- **Complex instruction set computer (CISC)**
 - Many different instructions with many different formats
 - But, only small subset encountered with Linux programs
 - Hard to match performance of Reduced Instruction Set Computers (RISC)
 - But, Intel has done just that!
 - In terms of speed. Less so for low power.

Intel x86 Evolution: Milestones

<i>Name</i>	<i>Date</i>	<i>Transistors</i>	<i>MHz</i>
■ 8086	1978	29K	5-10
<ul style="list-style-type: none">■ First 16-bit Intel processor. Basis for IBM PC & DOS■ 1MB address space			
■ 386	1985	275K	16-33
<ul style="list-style-type: none">■ First 32 bit Intel processor , referred to as IA32■ Added “flat addressing”, capable of running Unix			
■ Pentium 4E	2004	125M	2800-3800
<ul style="list-style-type: none">■ First 64-bit Intel x86 processor, referred to as x86-64			
■ Core 2	2006	291M	1060-3500
<ul style="list-style-type: none">■ First multi-core Intel processor			
■ Core i7	2008	731M	1700-3900
<ul style="list-style-type: none">■ Four cores			

Moore's Law

Aside Moore's law



If we plot the number of transistors in the different Intel processors versus the year of introduction, and use a logarithmic scale for the y-axis, we can see that the growth has been phenomenal. Fitting a line through the data, we see that the number of transistors increases at an annual rate of approximately 38%, meaning that the number of transistors doubles about every 26 months. This growth has been sustained over the multiple-decade history of x86 microprocessors.

2015 State of the Art

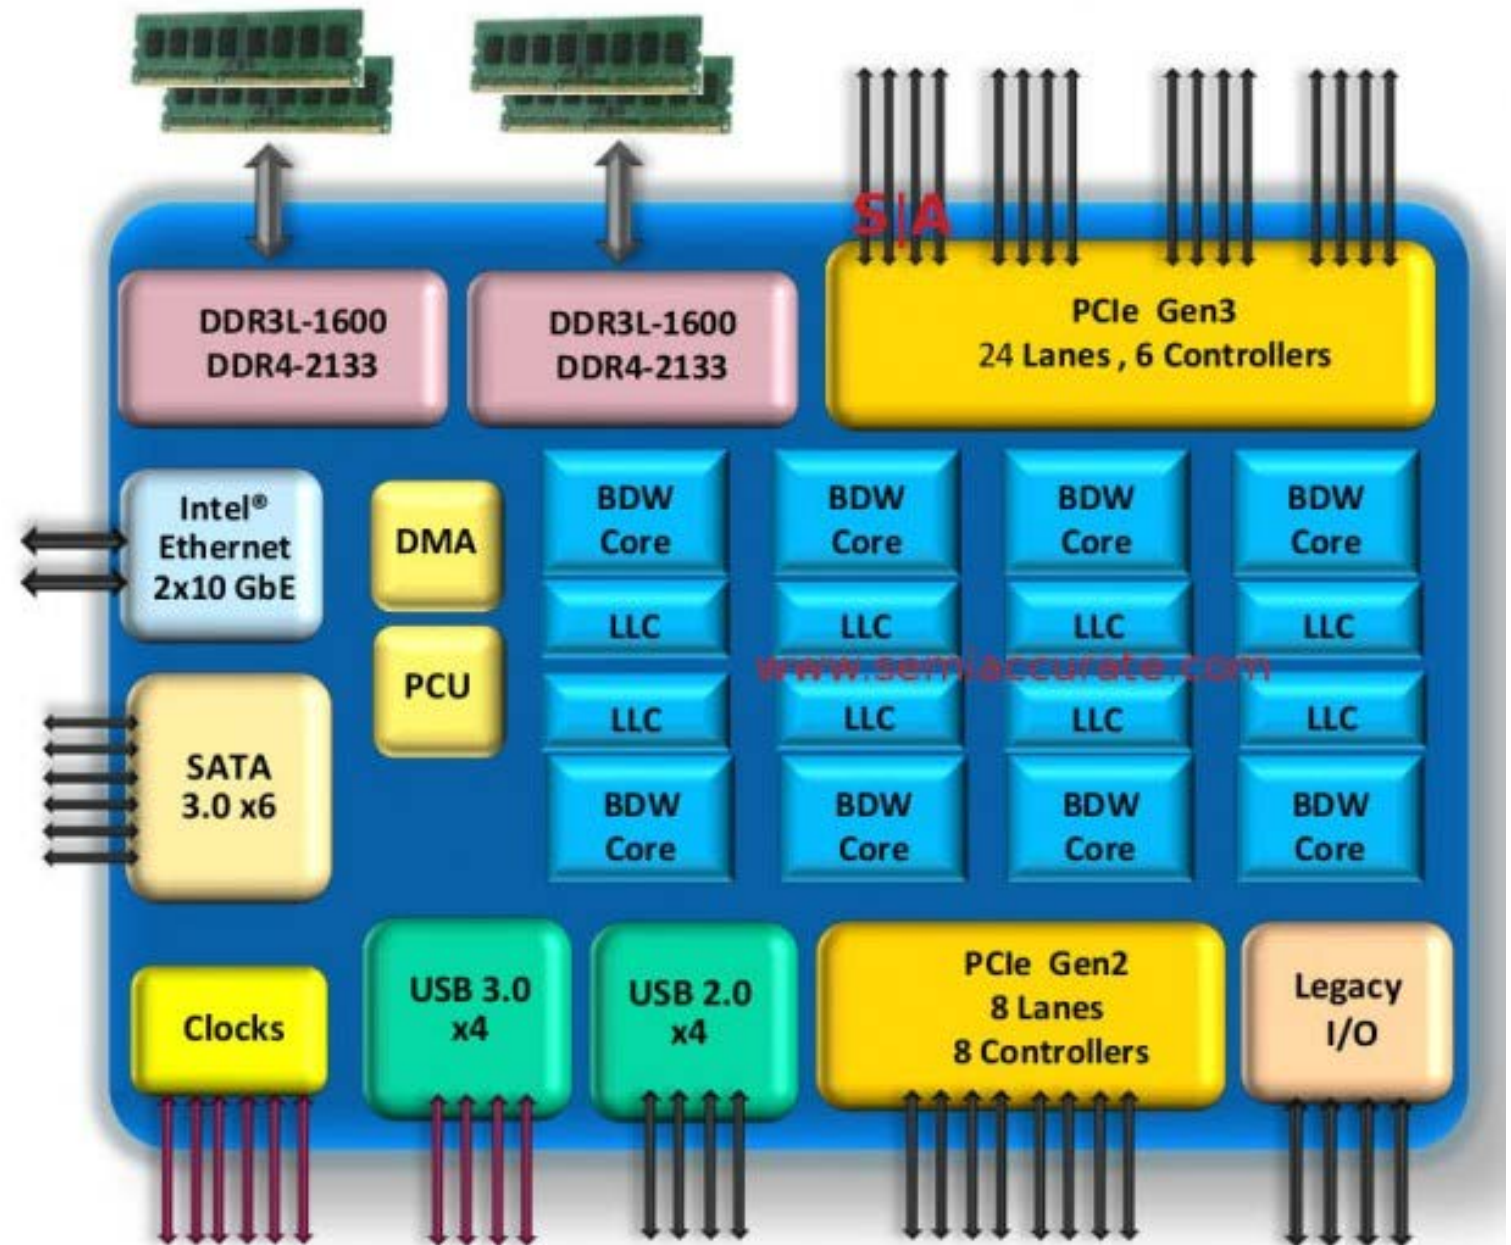
- Core i7 Broadwell 2015

■ Desktop Model

- 4 cores
- Integrated graphics
- 3.3-3.8 GHz
- 65W

■ Server Model

- 8 cores
- Integrated I/O
- 2-2.6 GHz
- 45W



Intel's 64-Bit History

- **2001: Intel Attempts Radical Shift from IA32 to IA64**
 - Totally different architecture (Itanium)
 - Executes IA32 code only as legacy
 - Performance disappointing
- **2003: AMD Steps in with Evolutionary Solution**
 - x86-64 (now called “AMD64”)
- **Intel Felt Obligated to Focus on IA64**
 - Hard to admit mistake or that AMD is better
- **2004: Intel Announces EM64T extension to IA32**
 - Extended Memory 64-bit Technology
 - Almost identical to x86-64!
- **All but low-end x86 processors support x86-64**
 - But, lots of code still runs in 32-bit mode

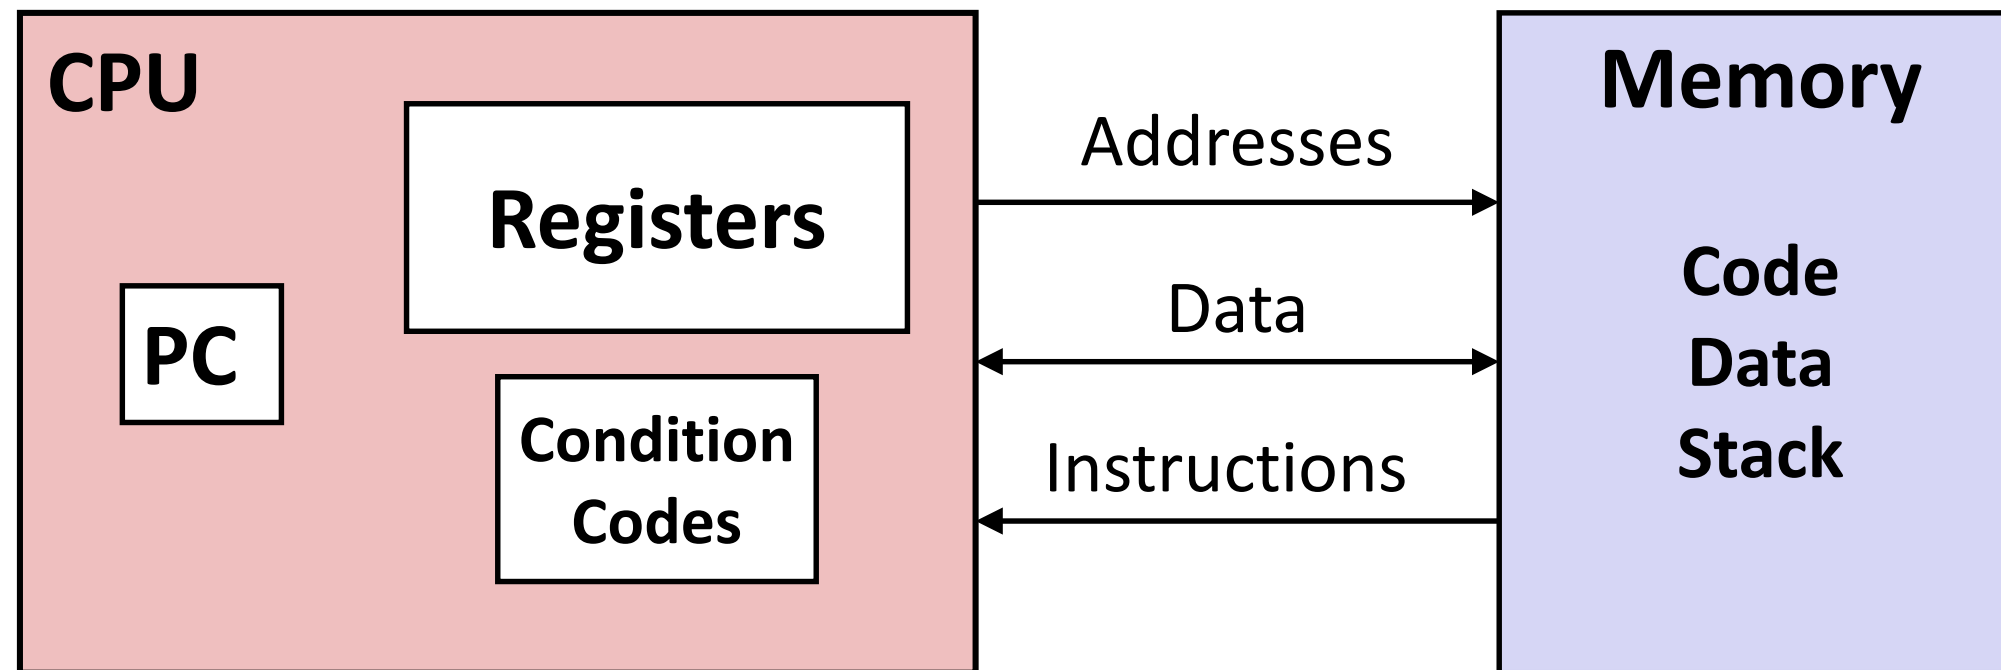
Today: Machine Programming I: Basics

- History of Intel processors and architectures
- **C, assembly, machine code**
- Assembly Basics: Registers, operands, move
- Arithmetic & logical operations

Definitions

- **Architecture:** (also ISA: instruction set architecture) The parts of a processor design that one needs to understand or write assembly/machine code.
 - Examples: instruction set specification, registers.
- **Code Forms:**
 - **Machine Code:** The byte-level programs that a processor executes
 - **Assembly Code:** A text representation of machine code
- **Example ISAs:**
 - Intel: x86, IA32, Itanium, x86-64
 - ARM: Used in almost all mobile phones

Assembly/Machine Code View



Programmer-Visible State

- **PC: Program counter**

- Address of next instruction
- Called “RIP” (x86-64)

- **Register file**

- Heavily used program data

- **Condition codes**

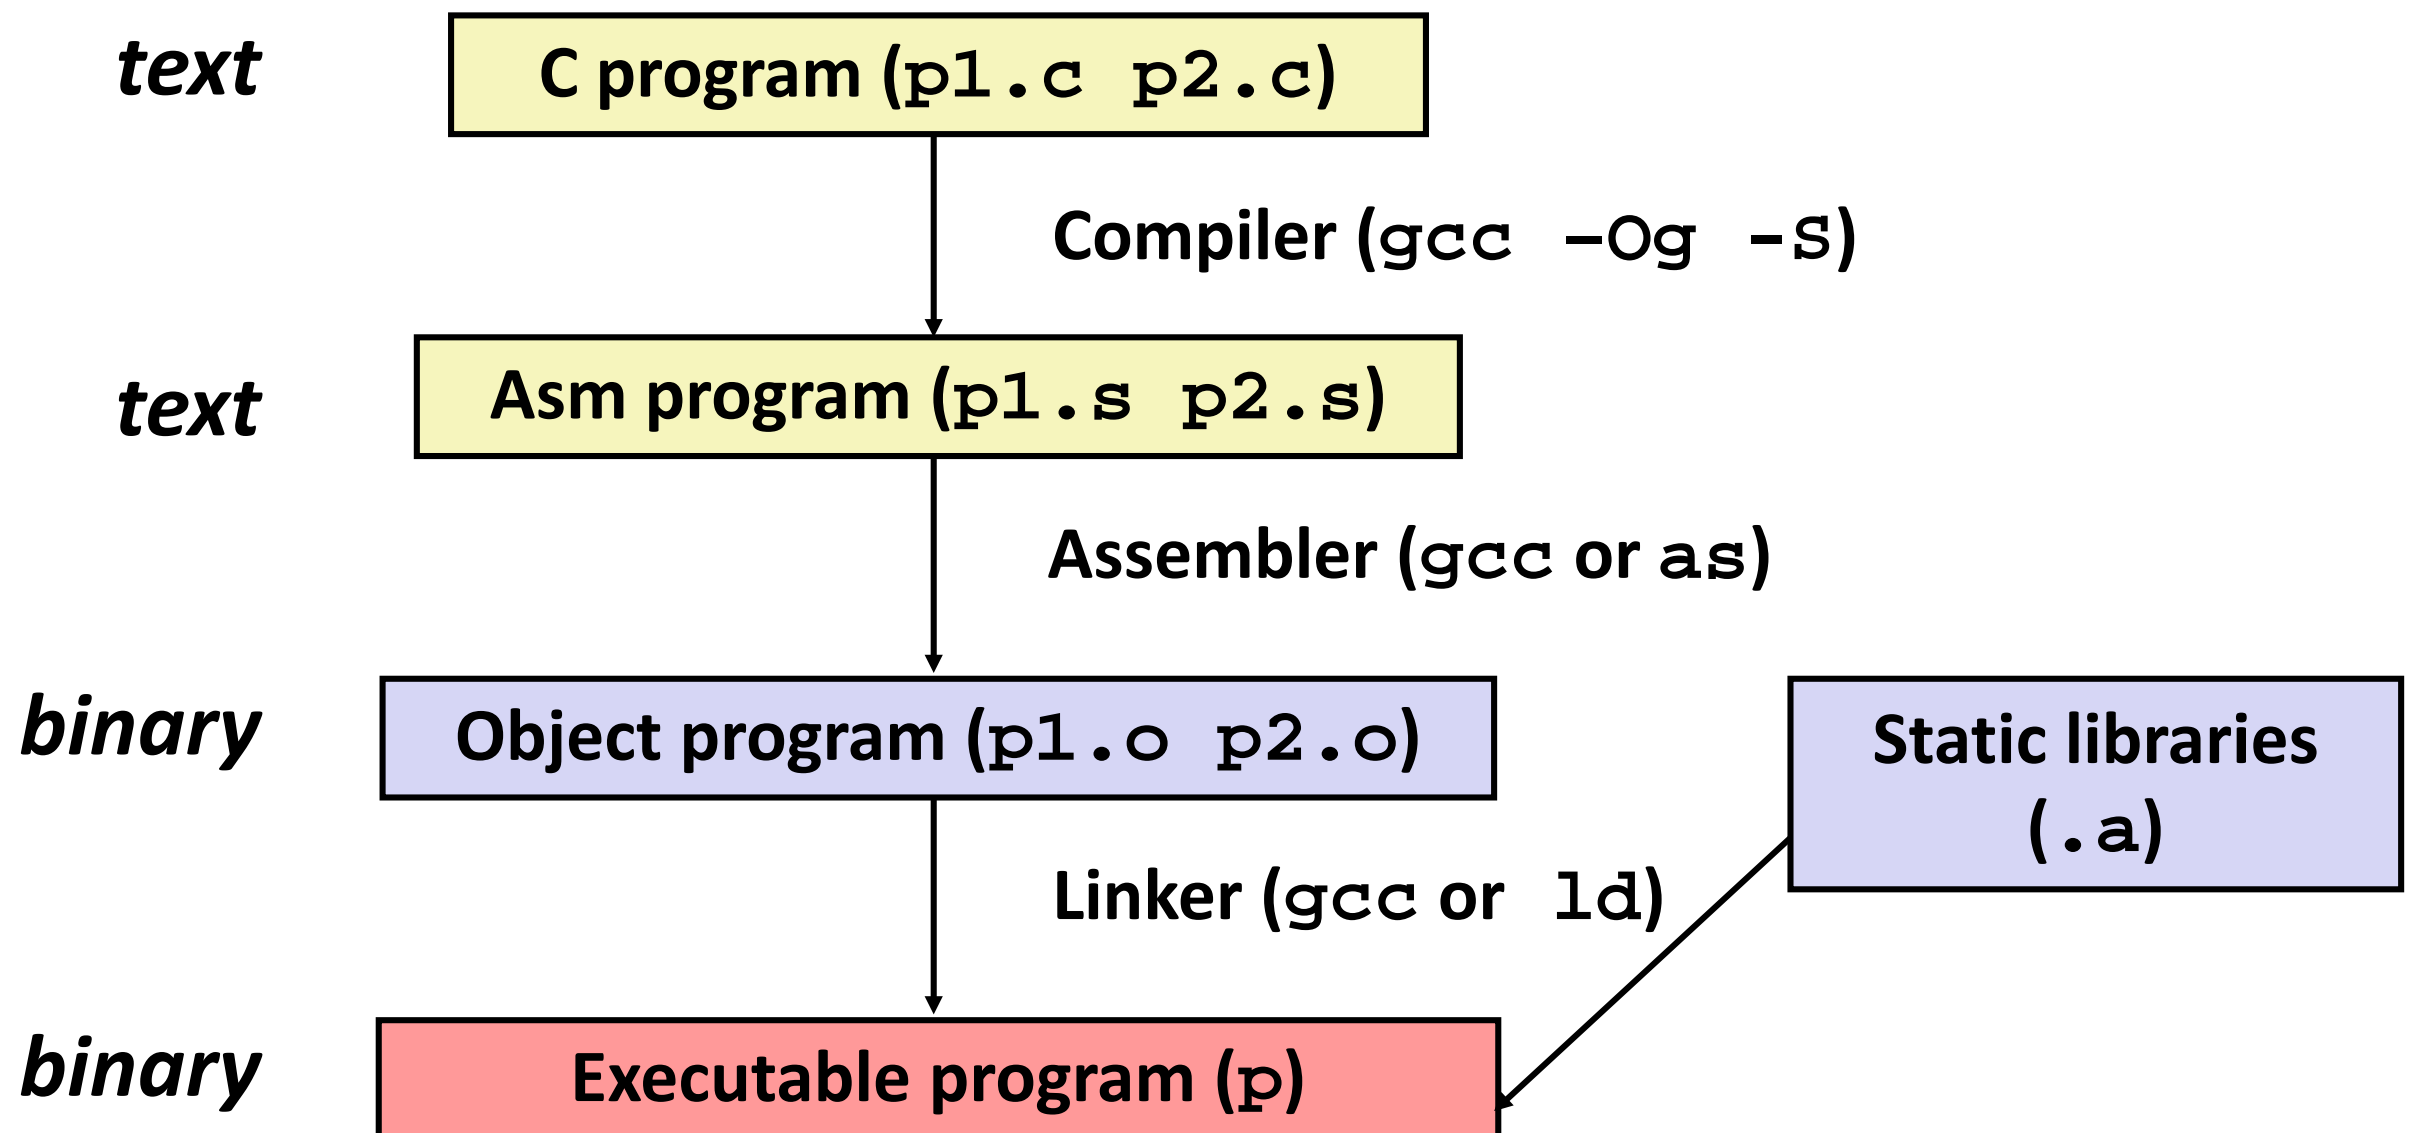
- Store status information about most recent arithmetic or logical operation
- Used for conditional branching

- **Memory**

- Byte addressable array
- Code and user data
- Stack to support procedures

Turning C into Object Code

- Code in files `p1.c` `p2.c`
- Compile with command: `gcc -Og p1.c p2.c -o p`
 - Use basic optimizations (`-Og`) [New to recent versions of GCC]
 - Put resulting binary in file `p`



Compiling Into Assembly

C Code (sum.c)

```
long plus(long x, long y);

void sumstore(long x, long y,
              long *dest)
{
    long t = plus(x, y);
    *dest = t;
}
```

Generated x86-64 Assembly

```
sumstore:
    pushq    %rbx
    movq     %rdx, %rbx
    call     plus
    movq     %rax, (%rbx)
    popq     %rbx
    ret
```

Obtain (on windows machine) with command

```
gcc -Og -fno-asynchronous-unwind-tables -S
sum.c
```

Produces file `sum.s`

Warning: Will get very different results on non-windows machines (Linux, Mac OS-X, ...) due to different versions of gcc and different compiler settings.

Assembly Characteristics: Data Types

- **“Integer” data of 1, 2, 4, or 8 bytes**
 - Data values
 - Addresses (untyped pointers)
- **Floating point data of 4, 8, or 10 bytes**
- **Code: Byte sequences encoding series of instructions**
- **No aggregate types such as arrays or structures**
 - Just contiguously allocated bytes in memory

Assembly Characteristics: Operations

- **Perform arithmetic function on register or memory data**
- **Transfer data between memory and register**
 - Load data from memory into register
 - Store register data into memory
- **Transfer control**
 - Unconditional jumps to/from procedures
 - Conditional branches

Object Code

Code for `sumstore`

`0x0400595:`

`0x53`

`0x48`

`0x89`

`0xd3`

`0xe8`

`0xf2`

`0xff`

`0xff`

`0xff`

`0x48`

`0x89`

`0x03`

`0x5b`

`0xc3`

- Total of 14 bytes
- Each instruction 1, 3, or 5 bytes
- Starts at address `0x0400595`

■ Assembler

- Translates `.s` into `.o`
- Binary encoding of each instruction
- Nearly-complete image of executable code
- Missing linkages between code in different files

■ Linker

- Resolves references between files
- Combines with static run-time libraries
 - E.g., code for `malloc`, `printf`
- Some libraries are *dynamically linked*
 - Linking occurs when program begins execution

Machine Instruction Example

```
*dest = t;
```

```
movq %rax, (%rbx)
```

```
0x40059e:  48 89 03
```

■ C Code

- Store value `t` where designated by `dest`

■ Assembly

- Move 8-byte value to memory
 - Quad words in x86-64 parlance
- Operands:
 - `t`: Register `%rax`
 - `dest`: Register `%rbx`
 - `*dest`: Memory `M[%rbx]`

■ Object Code

- 3-byte instruction
- Stored at address `0x40059e`

Disassembling Object Code

Disassembled

```
00000000000400595 <sumstore>:
400595:  53                push    %rbx
400596:  48 89 d3          mov     %rdx,%rbx
400599:  e8 f2 ff ff ff    callq   400590 <plus>
40059e:  48 89 03          mov     %rax, (%rbx)
4005a1:  5b                pop     %rbx
4005a2:  c3                retq
```

■ Disassembler

`objdump -d sum`

- Useful tool for examining object code
- Analyzes bit pattern of series of instructions
- Produces approximate rendition of assembly code
- Can be run on either a `.out` (complete executable) or `.o` file

Alternate Disassembly

Object

0x0400595:

0x53

0x48

0x89

0xd3

0xe8

0xf2

0xff

0xff

0xff

0x48

0x89

0x03

0x5b

0xc3

Disassembled

Dump of assembler code for function sumstore:

0x000000000000400595 <+0>: push %rbx

0x000000000000400596 <+1>: mov %rdx,%rbx

0x000000000000400599 <+4>: callq 0x400590 <plus>

0x00000000000040059e <+9>: mov %rax, (%rbx)

0x0000000000004005a1 <+12>: pop %rbx

0x0000000000004005a2 <+13>: retq

■ Within gdb Debugger

`gdb sum`

`disassemble sumstore`

■ Disassemble procedure

`x/14xb sumstore`

■ Examine the 14 bytes starting at sumstore

What Can be Disassembled?

```
% objdump -d WINWORD.EXE
```

```
WINWORD.EXE:      file format pei-i386
```

```
No symbols in "WINWORD.EXE".
```

```
Disassembly of section .text:
```

```
30001000 <.text>:
```

```
30001000:
```

```
30001001:
```

```
30001003:
```

```
30001005:
```

```
3000100a:
```

**Reverse engineering forbidden by
Microsoft End User License Agreement**

- Anything that can be interpreted as executable code
- Disassembler examines bytes and reconstructs assembly source

Today: Machine Programming I: Basics

- History of Intel processors and architectures
- C, assembly, machine code
- **Assembly Basics: Registers, operands, move**

Some History: IA32 Registers

Origin
(mostly obsolete)

general purpose	%eax	%ax	%ah	%al	accumulate
	%ecx	%cx	%ch	%cl	counter
	%edx	%dx	%dh	%dl	data
	%ebx	%bx	%bh	%bl	base
	%esi	%si			source index
	%edi	%di			destination index
	%esp	%sp			stack pointer
	%ebp	%bp			base pointer
16-bit virtual registers (backwards compatibility)					

x86 (32-bit)

- General purpose registers:
EAX, EBX, ECX, EDX, ESI, EDI
(though some have historical uses: **ECX** is used for counters, **ESI & EDI** are source and destination for string operations)
- Special-purpose registers:
EBP (base pointer), **ESP** (stack pointer)
- Flags register: **EFLAGS**
- Segment registers: **CS, SS, DS, ES, FS, GS**
- Instruction pointer/program counter: **EIP**
- Control registers: **CR0-CR4**

x86 Instructions

- Instructions are *variable length*
- Conventions:
 - regN means a register of N bits
 - mem means memory
 - constN means a constant of N bits

Note: x86 is *BIG*

- **When they stopped printing paper manuals (c. 2010), occupied Intel manual occupied 5 volumes**
- **Two volumes (1300 pages as of June 2015) dedicated just to describing instructions**
- **We will only cover a minimal subset of x86 here; if you have questions about a particular instruction the Intel manuals are actually pretty good**



Intel® 64 and IA-32 Architectures Software Developer's Manual
VOLUME 3B: System Programming Guide Part 2

3B



Intel® 64 and IA-32 Architectures Software Developer's Manual
VOLUME 3A: System Programming Guide Part 1

3A



Intel® 64 and IA-32 Architectures Software Developer's Manual
VOLUME 2B: Instruction Set Reference, N-Z

2B



Intel® 64 and IA-32 Architectures Software Developer's Manual
VOLUME 2A: Instruction Set Reference, A-M

2A



Intel® 64 and IA-32 Architectures Software Developer's Manual
VOLUME 1: Basic Architecture

1

x86 Assembly Syntax

- **Two "flavors":**

- Intel syntax (what's in the processor manual)
- AT&T syntax (what GNU as uses)

- **The difference between them:**

- Intel: `mov eax, DWORD PTR [ebx]`
- AT&T: `movl (%ebx), %eax`

- **GNU as can be made to use Intel syntax instead with `.intel_syntax noprefix`**

Data Movement

- **mov <reg>,<reg>**
- **mov <reg>,<mem>**
- **mov <mem>,<reg>**
- **mov <reg>,<const>**
- **mov <mem>,<const>**

Data Movement

■ Examples (AT&T syntax)

- Source, dest
- `movl %eax, %edx` \Rightarrow `edx = eax`
- `movl $0x123, %ebx` \Rightarrow `ebx = 0x123`
- `movl 0x123, %ebx` \Rightarrow `ebx = *(int *)0x123`
- `movl (%ebx), %edx` \Rightarrow `edx = *(int *)ebx`
- `movl 4(%ebx), %edx` \Rightarrow `edx = *(int *)(ebx + 4)`

Segmentation

- Original 8086 processor had 20 *address lines* (i.e., bits for specifying a memory address)
- But only had 16-bit registers
- Segment registers provide a way to get the extra 4 bits
- Segments also have permissions associated with them (not discussed here)
- Modern OSes make almost no use of segments (and x86_64 gets rid of them in 64-bit mode)

Segmentation

- **cs: *code segment*** (modifies code fetch)
- **ds: *data segment*** (modifies data access)
- **ss: *stack segment*** (modifies ESP & EBP)
- **es: *extra segment*** (modifies string ops)
- **So:**
 - `movl (%ebp), %ebx` implicitly uses the 4 bits in the ss as a prefix to %ebp

Push/Pop

- Recall: ESP is stack pointer
- On x86 the stack grows *downward* – putting something on the stack decrements the stack pointer
- Push subtracts 4 from ESP, then writes to the memory at address ESP

Arithmetic


- add, sub, inc, dec
- imul, idiv
- and, or, xor
- not
- neg
- shl, shr

Logical

- **test %x, %y** – equivalent to **and %x, %y** but does not modify %y
- **cmp %x, %y** – equivalent to **sub %x, %y** but does not modify %y
- **So if they don't modify the destination operand, what do they do?**

EFLAGS

- Arithmetic and logical instructions set bits in EFLAGS as a side effect



Reserved bit positions. DO NOT USE.
 Always set to values previously read.

EFLAGS

- Conditional branches depend on the values in EFLAGS
- Also some specific instructions, e.g. `cmov`, will be executed only when flags are satisfied
- Other architectures take this even further: on ARM, almost *any* instruction can be made conditional

Unconditional Branch

■ **jmp: unconditional jump**

- jmp <const> – jump to constant address
- jmp <mem> – dereference mem and jump
- jmp <reg> - jump to address in register

Conditional Branches

Instruction	Description	signed-ness	Flags
JO	Jump if overflow		OF = 1
JNO	Jump if not overflow		OF = 0
JS	Jump if sign		SF = 1
JNS	Jump if not sign		SF = 0
JE JZ	Jump if equal Jump if zero		ZF = 1
JNE JNZ	Jump if not equal Jump if not zero		ZF = 0
JB JNAE JC	Jump if below Jump if not above or equal Jump if carry	unsigned	CF = 1
JNB JAE JNC	Jump if not below Jump if above or equal Jump if not carry	unsigned	CF = 0
JBE JNA	Jump if below or equal Jump if not above	unsigned	CF = 1 or ZF = 1
JA JNBE	Jump if above Jump if not below or equal	unsigned	CF = 0 and ZF = 0
JL JNGE	Jump if less	signed	SF <> OF

System

- Many instructions control system behavior
- **int** – executes a *software interrupt* (traps into the kernel)
- **iret** – return from an interrupt
- **sysenter/sysret** – newer way to enter/leave kernel mode

Fill in the blanks

Address	Value	Register	Value
0x100	0xFF	%rax	0x100
0x104	0xAB	%rcx	0x1
0x108	0x13	%rdx	0x3
0x10C	0x11		

Operand	Value	Comment
%rax		
0x104		
\$0x108		
(%rax)		
4(%rax)		