

# Lecture 5: Assembly Boot

Professor G. Sandoval

Some slides adapted by G. Sandoval for CS3224, from Tanenbaum & Bo, Modern Operating Systems: 4th ed., (c) 2013 Prentice-Hall, Inc. All rights reserved.  
Also some Slides by Brendan Dolan-Gavitt and Bryant and O'Hallaron, Computer Systems: A programmer's Perspective, Third Edition

# Today: Machine Programming Basics

- **Assembly**

- More examples
- Function calls

- **GDB**

- **Booting the OS**

- BIOS

# Fill in the blanks

Address	Value	Register	Value
0x100	0xFF	%rax	0x100
0x104	0xAB	%rcx	0x1
0x108	0x13	%rdx	0x3
0x10C	0x11		

Operand	Value	Comment
%rax		
0x104		
\$0x108		
(%rax)		
4(%rax)		

# Fill in the blanks

Address	Value	Register	Value
0x100	0xFF	%rax	0x100
0x104	0xAB	%rcx	0x1
0x108	0x13	%rdx	0x3
0x10C	0x11		

Operand	Value	Comment
%rax	0x100	Register
0x104	0xAB	Absolute Address
\$0x108	0x108	Immediate
(%rax)	0xFF	Address 0x100
4(%rax)	0xAB	Address 0x104

# x64 Assembly: Jumps

Instruction	Effect
<code>jmp</code>	Always jump
<code>je / jz</code>	Jump if eq / zero
<code>jne / jnz</code>	Jump if !eq / !zero
<code>jg</code>	Jump if greater
<code>jge</code>	Jump if greater / eq
<code>jl</code>	Jump if less
<code>jle</code>	Jump if less / eq

# x64 Assembly: A Quick Drill

```
cmp $0x15213, %r12  
jge deadbeef
```

If \_\_\_\_\_, jump to addr  
0xdeadbeef

```
cmp %rax, %rdi  
jne 15213b
```

If \_\_\_\_\_, jump to addr  
0x15213b

```
test %r8, %r8  
jnz (%rsi)
```

If \_\_\_\_\_, jump to \_\_\_\_\_.

# x64 Assembly: A Quick Drill

<code>cmp \$0x15213, %r12</code>	<code>if %r12 &gt;= 0x15213,</code>
<code>jge deadbeef</code>	<code>jump to 0xdeadbeef</code>

```
cmp %rax, %rdi
jne 15213b
```

```
test %r8, %r8
jnz (%rsi)
```

# x64 Assembly: A Quick Drill

```
cmp $0x15213, %r12  
jge deadbeef
```

```
cmp %rax, %rdi  
jne 15213b
```

If the value of `%rdi` is not equal to value of `%rax`, jump to `0x15213b`.

```
test %r8, %r8  
jnz (%rsi)
```



# x64 Assembly: A Quick Drill

```
cmp $0x15213, %r12  
jge deadbeef
```

```
cmp %rax, %rdi  
jne 15213b
```

```
test %r8, %r8  
jnz (%rsi)
```

If `%r8 & %r8` is not zero,  
jump to the address stored in  
`%rsi`.

# Function Calls

- Supports creation of functions
- `call <f>`: push the address of next instruction on the stack and jump to `f`

- Example:

```
1 1000 call <f>           // pushes 5 and jumps to f
2 1005 movl %eax, %ebx
```

- `ret`: pop return address off the stack and jump to it

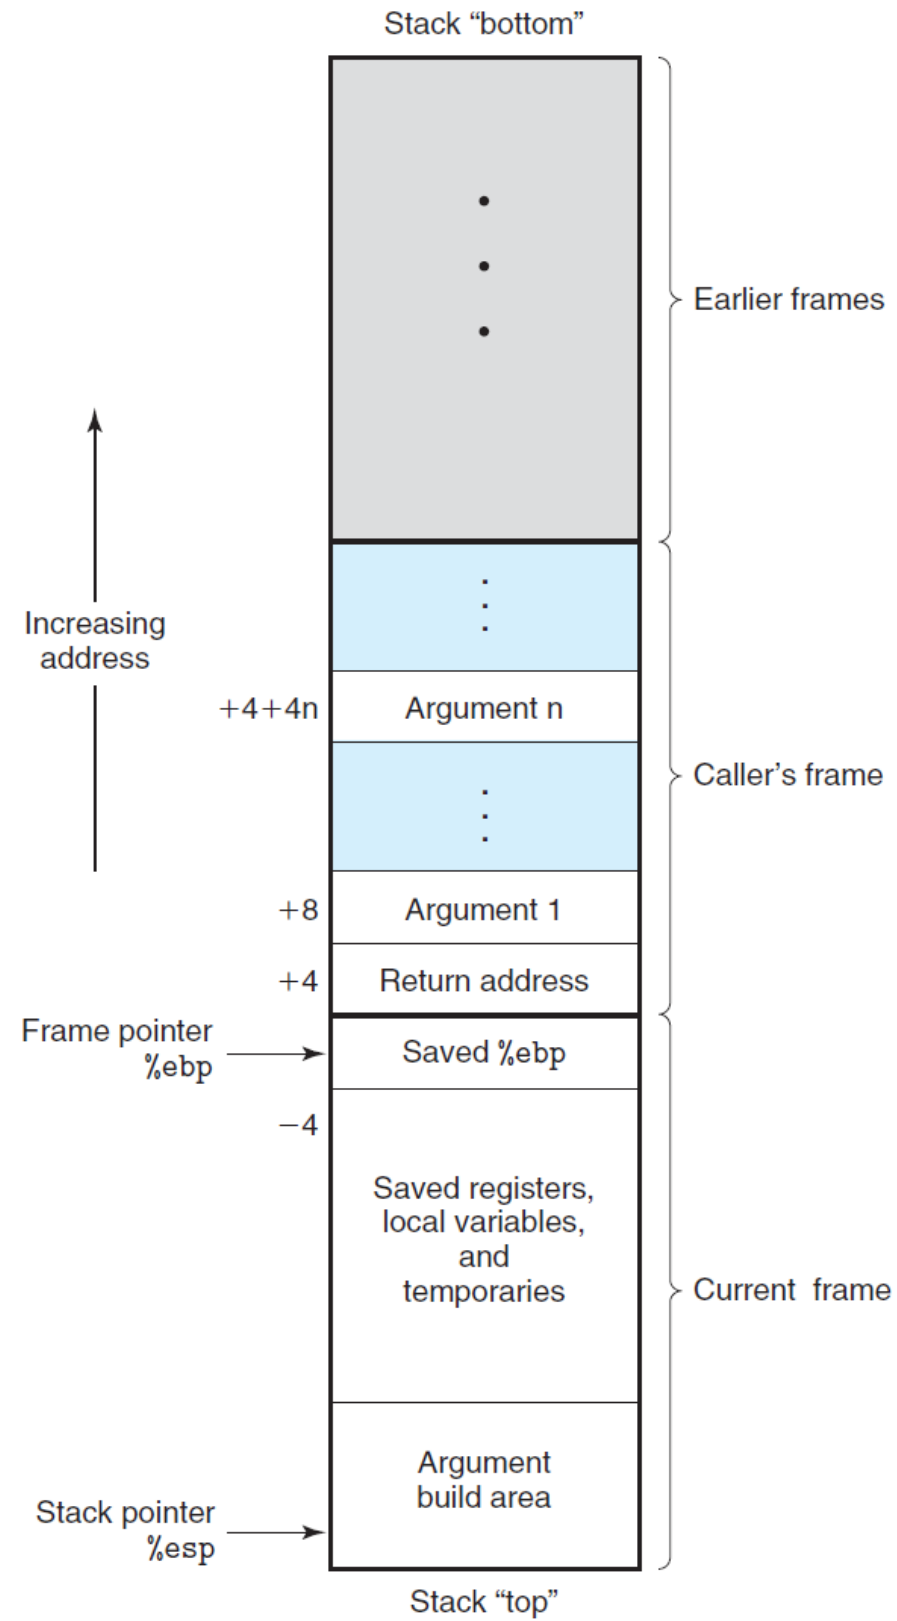
# Stack Frames

- **IA32 programs use the stack to support function calls.**
- **The machine uses the stack to:**
  - Pass arguments
  - store return information
  - Save registers
  - Local Storage
- **Therefore to use arguments and return information we need to deal with the stack.**
- **The portion of memory used in the stack for a single procedure call is called a stack frame.**

# Stack Frames

- The topmost stack frame is delimited by two register pointers:
  - **%ebp** is the frame pointer (base)
  - **%esp** is the stack pointer

# Stack Frame



# Function Prologue & Epilogue

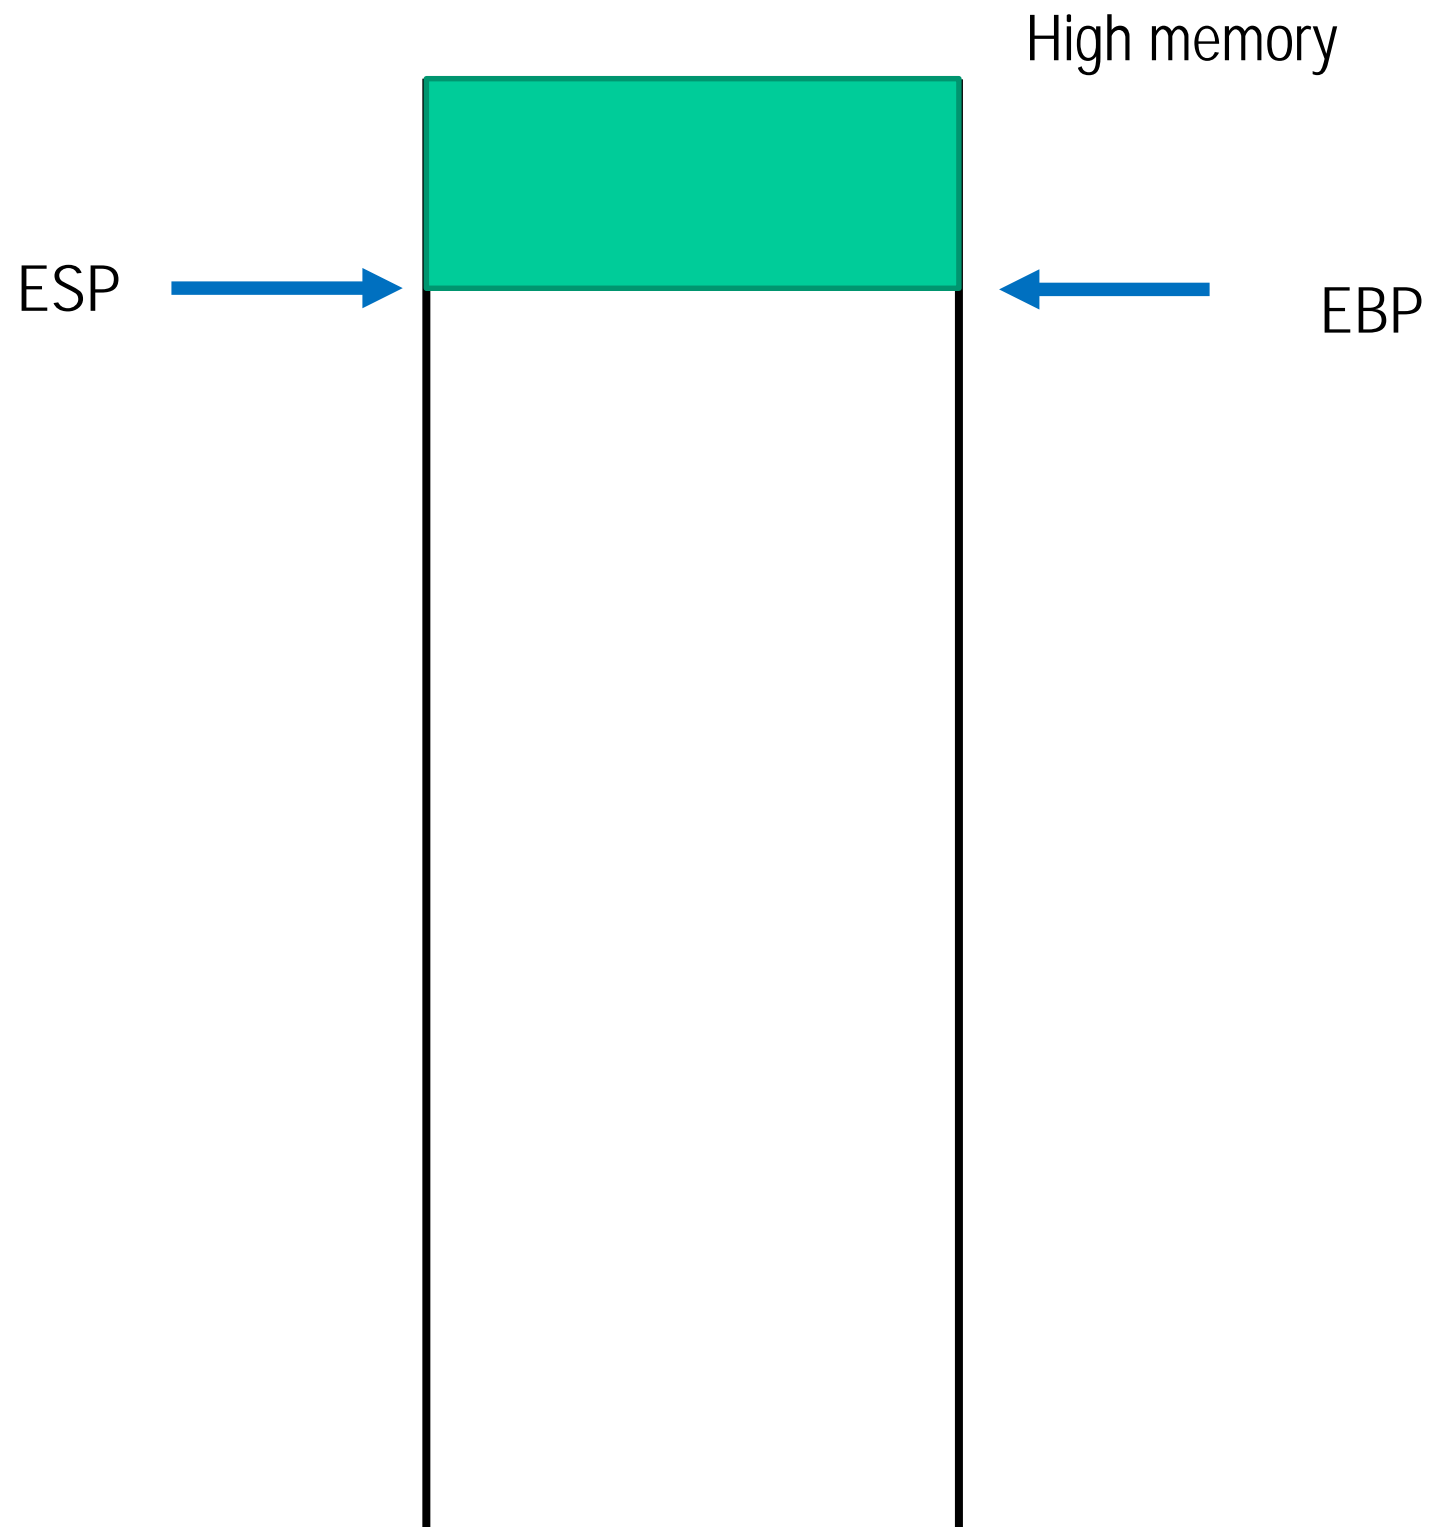
- The need to set up a stack frame gives rise to a common function prologue:

```
push %ebp           // save old value of ebp
mov %esp, %ebp      // set base p
sub $X, %esp        // reserve space for locals
```

- At the end, we can clean up with just the epilogue

```
add %ebp, %esp      // restore stack pointer
pop %ebp            // restore base pointer
ret                 // return pops return address
```

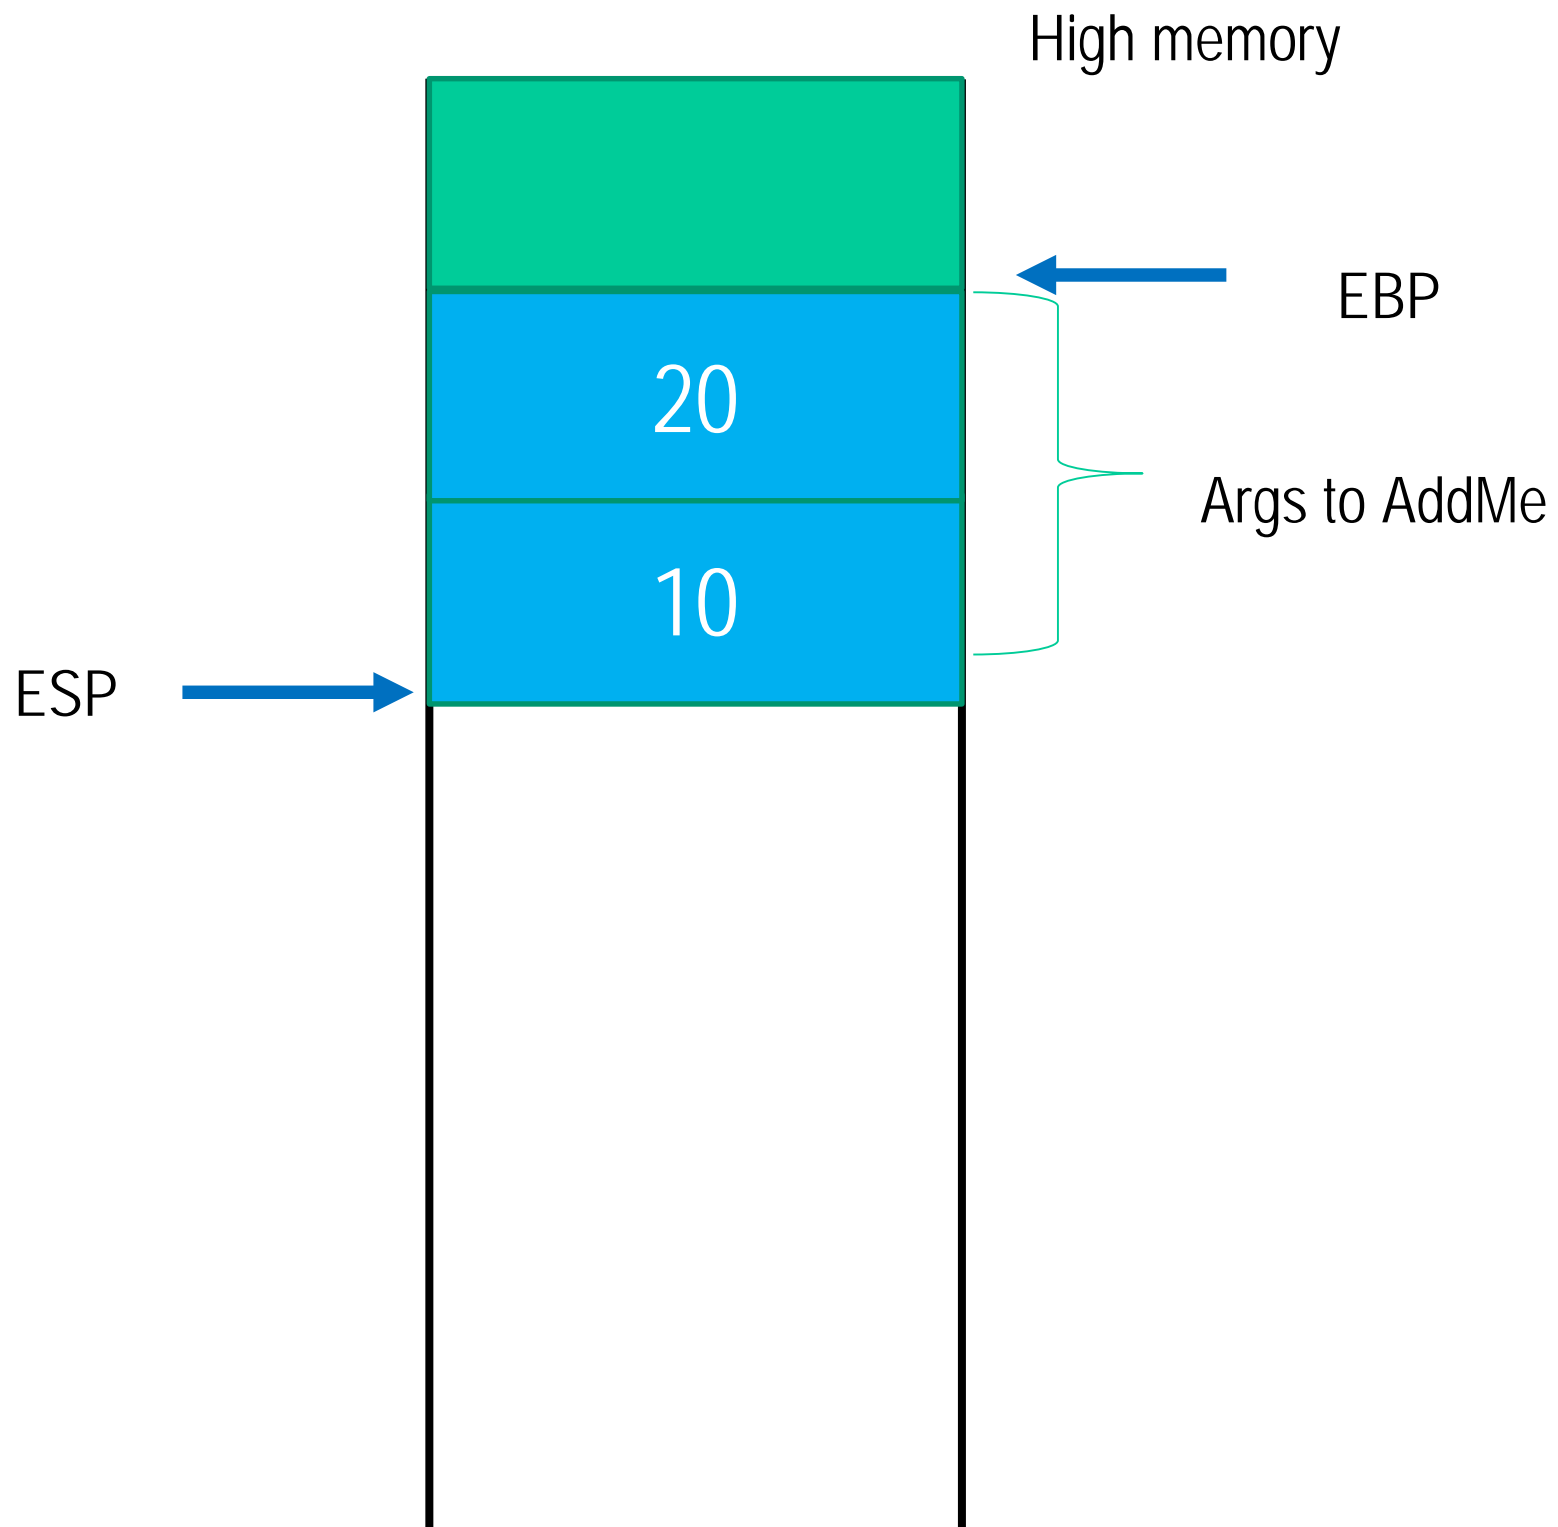
# Walking the Call Stack



```
int AddMe(int a, int b)
{
    int c;
    c = a + b;
    return c;
}

main()
{
    AddMe(10, 20);
    print...
}
```

# Walking the Call Stack

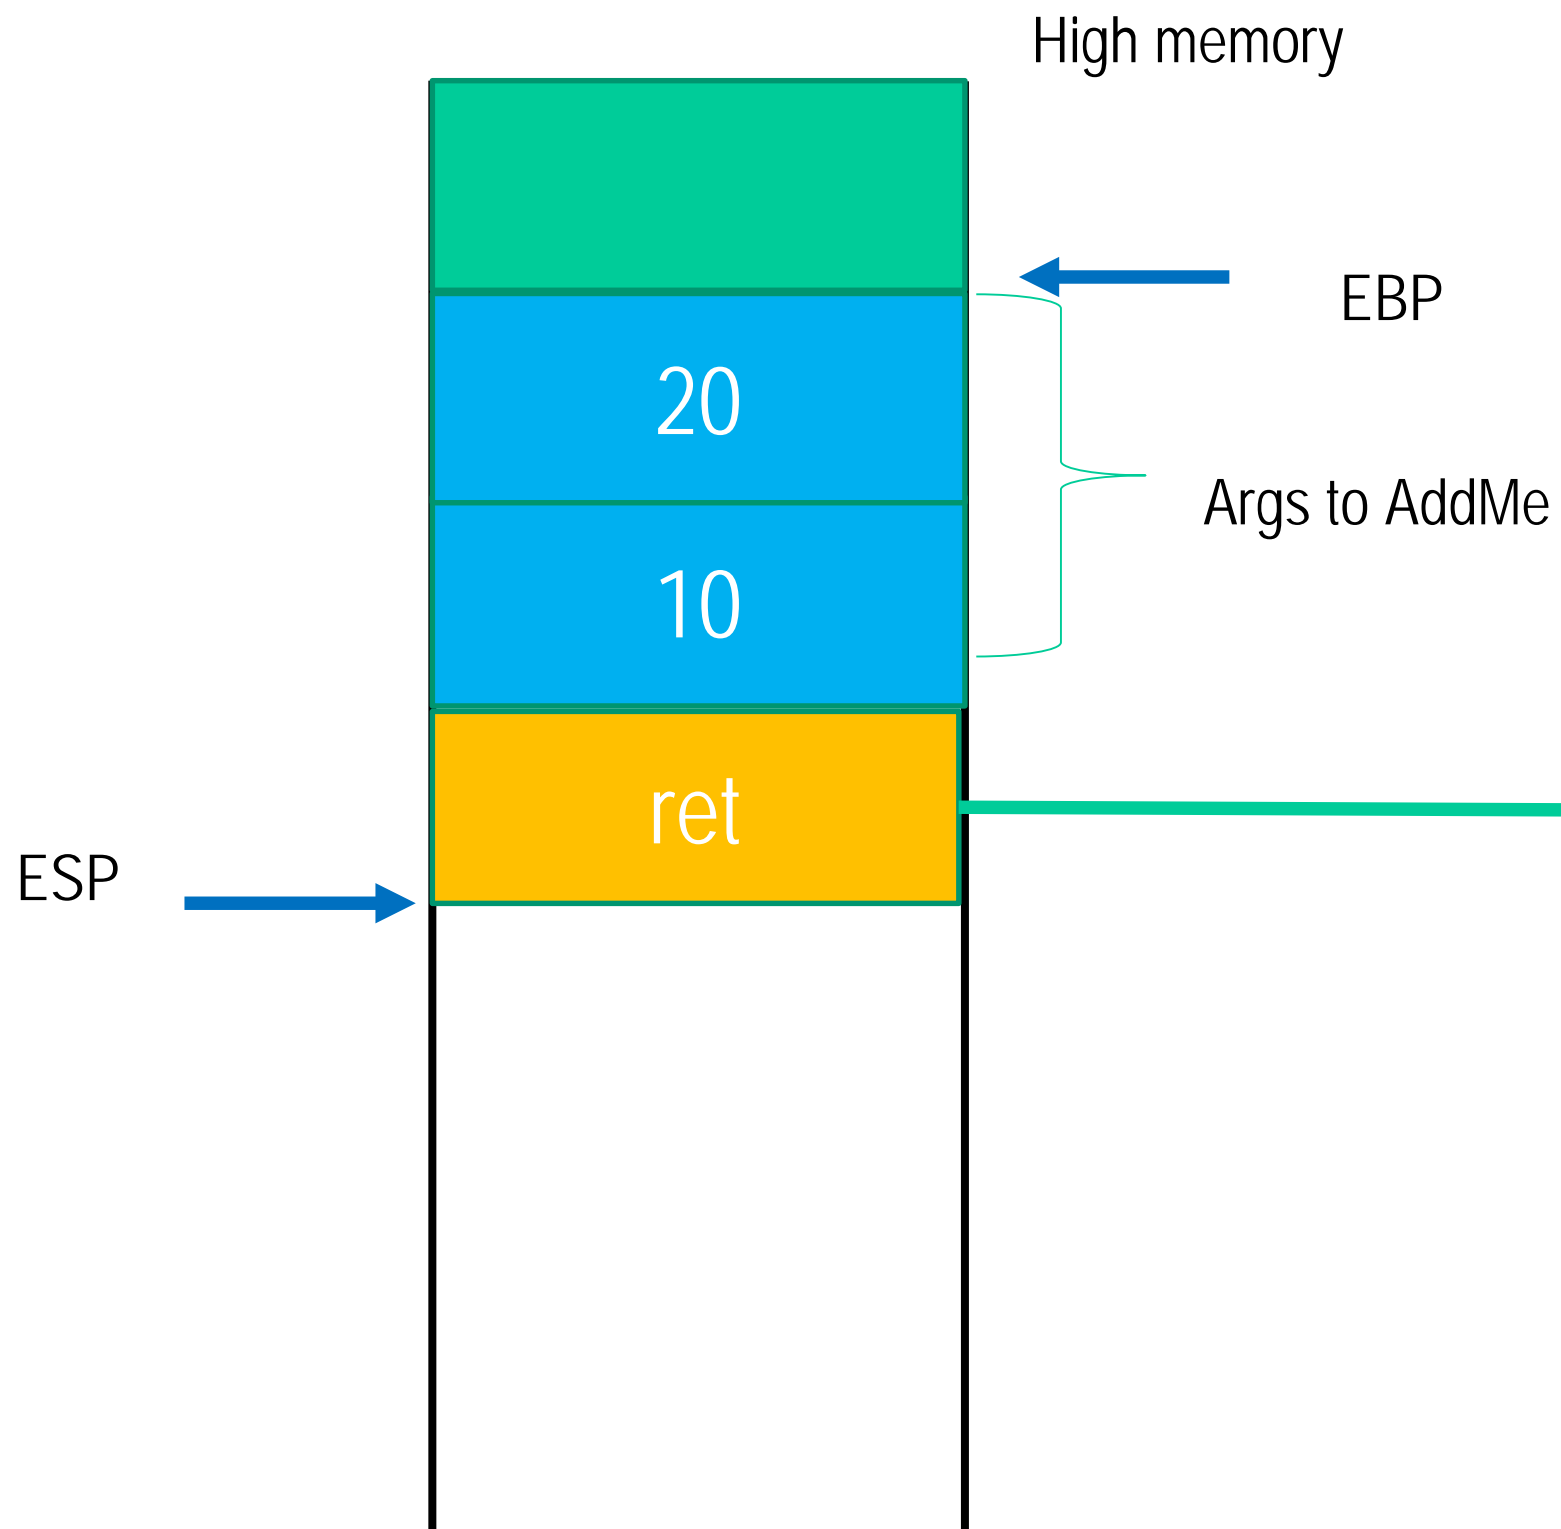


```
int AddMe(int a, int b)
{
    int c;
    c = a + b;
    return c;
}
```

```
main()
{
    AddMe(10, 20);
    print...
}
```



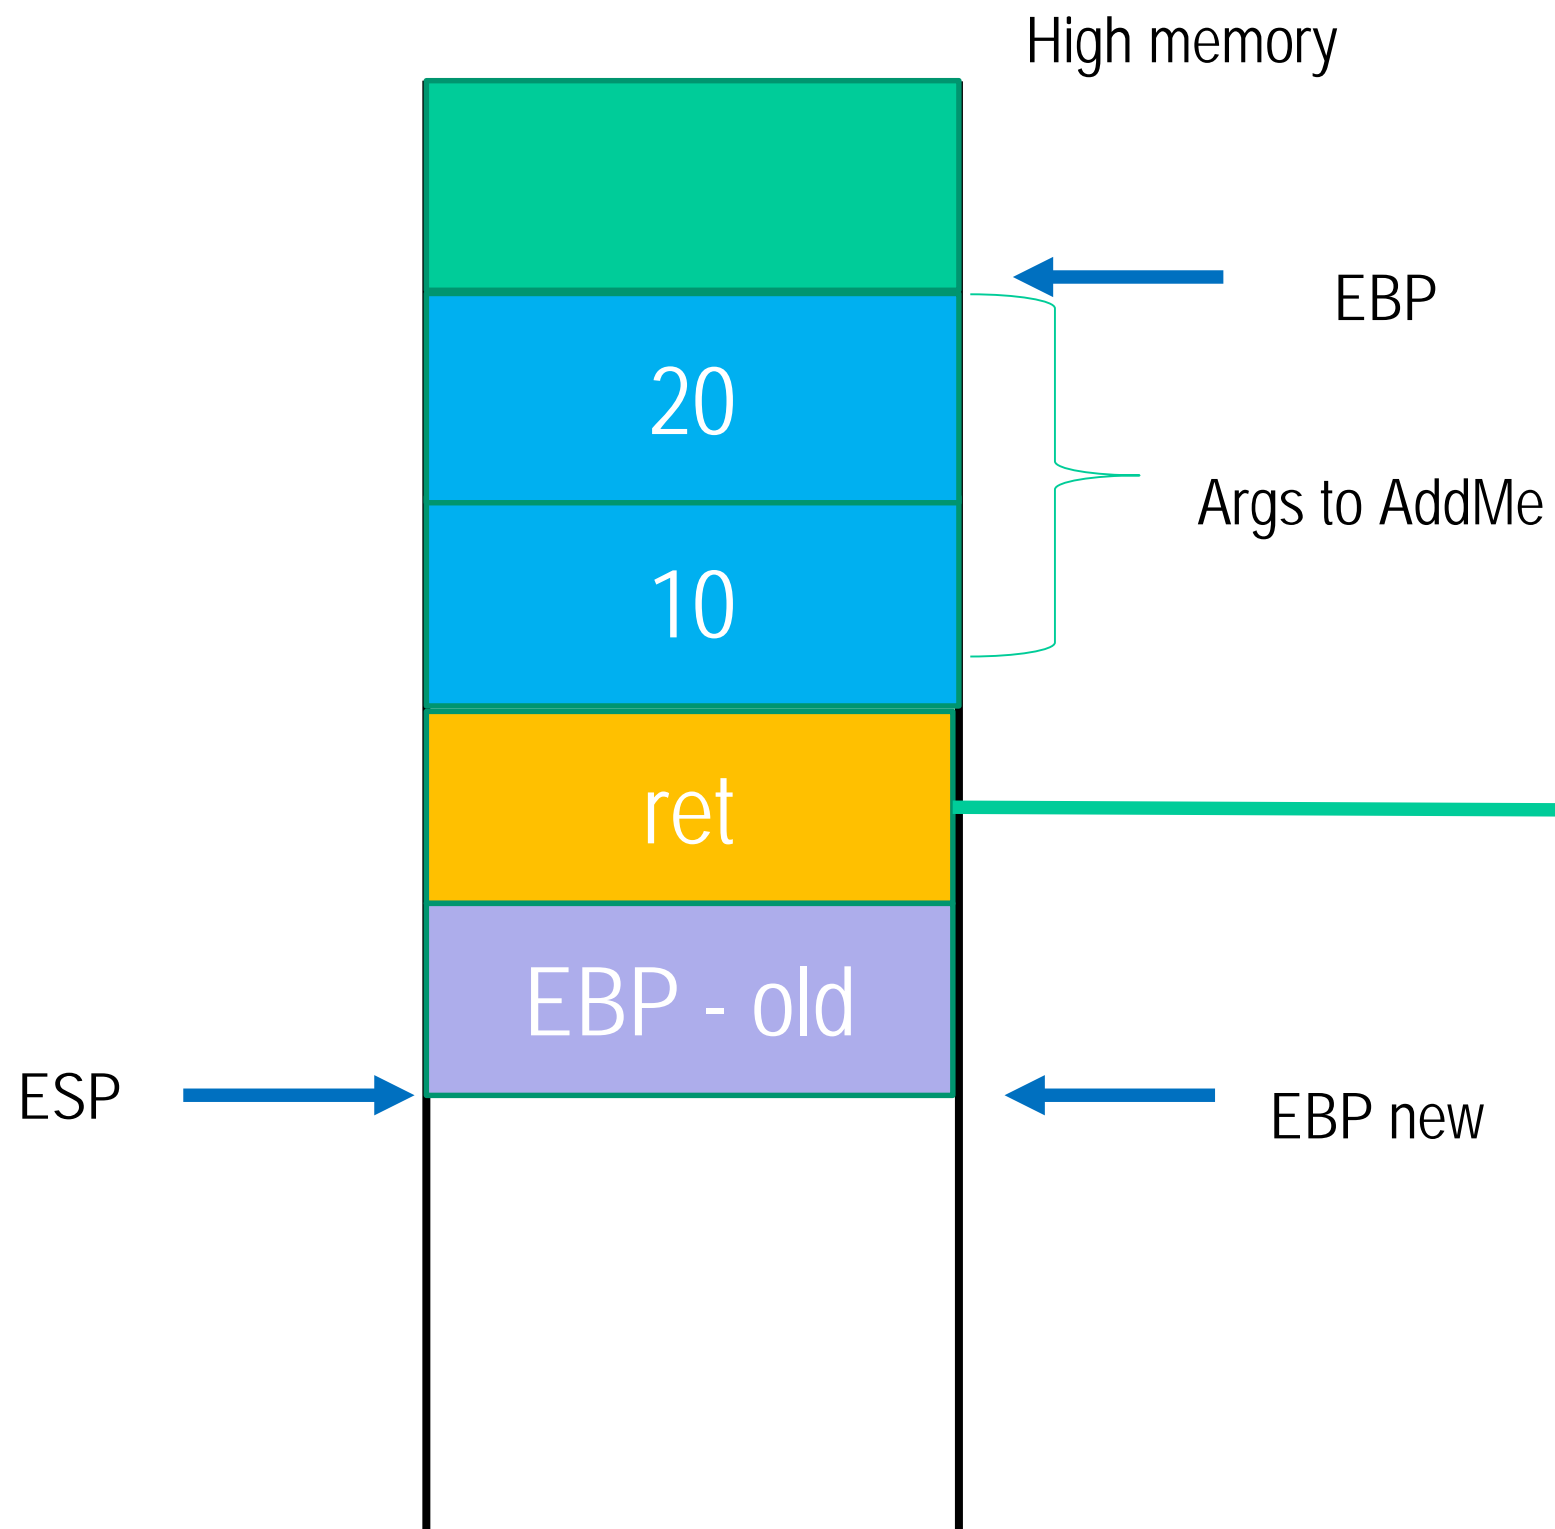
# Walking the Call Stack



```
int AddMe(int a, int b)
{
    int c;
    c = a + b;
    return c;
}

main()
{
    AddMe(10, 20);
    print...
}
```

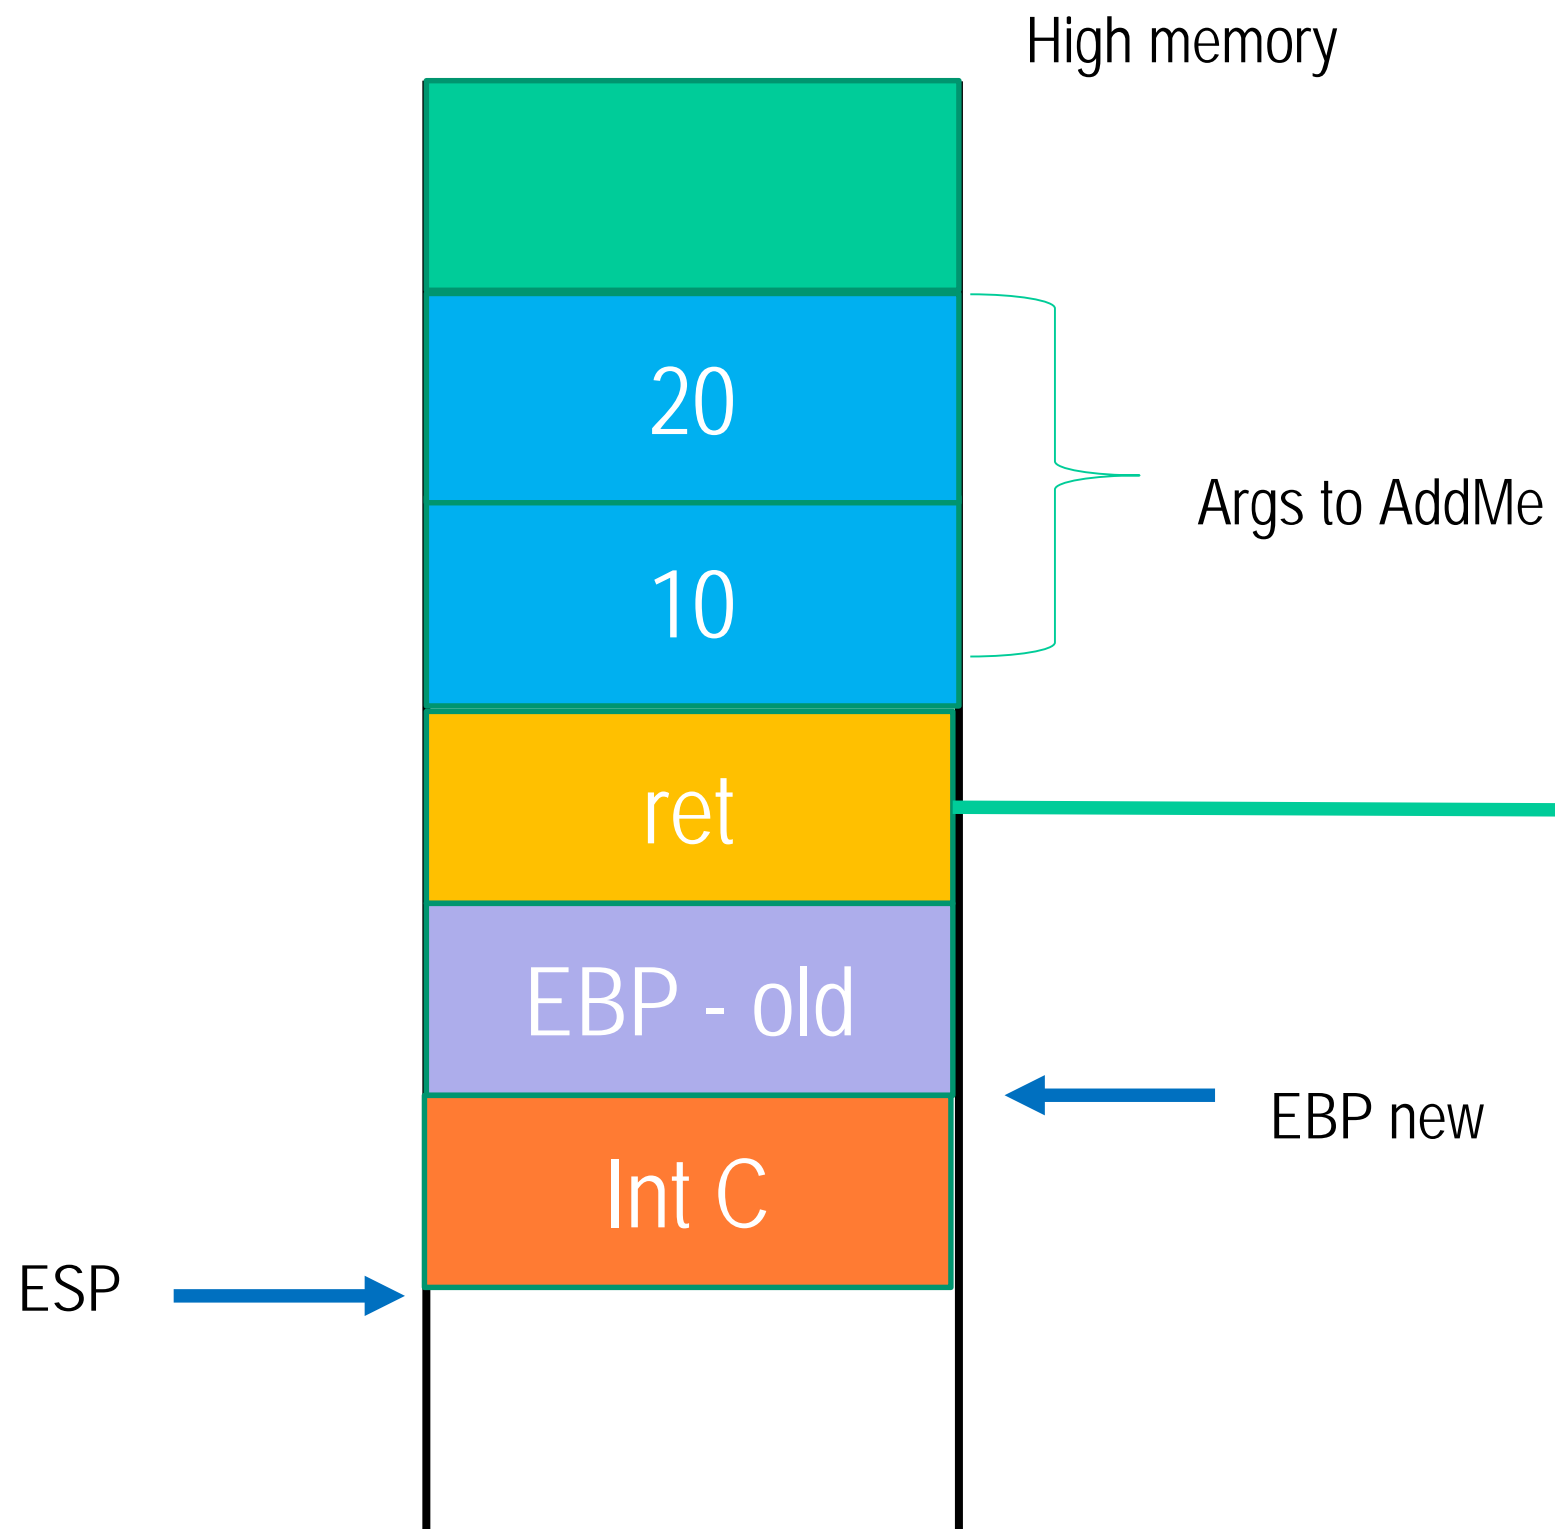
# Walking the Call Stack



```
int AddMe(int a, int b)
{
    int c;
    c = a + b;
    return c;
}

main()
{
    AddMe(10, 20);
    print...
}
```


# Walking the Call Stack



```
int AddMe(int a, int b)
{
    int c;
    c = a + b;
    return c;
}

main()
{
    AddMe(10, 20);
    print...
}
```

# To reference data on the stack



# Calling Conventions

- ***A calling convention* is a system-defined convention that tells you how to call a function**
- **In general a calling convention will:**
  - Say how to pass arguments to the function
  - Say who is responsible for cleaning up arguments
  - Specify where the return value will be placed
  - Say which registers will be preserved by the callee and which might be trashed

# gcc Calling Convention (32-bit)

- Return value: %eax
- Arguments are passed on the stack:

- f(1, 2, 3)
- push 3
- push 2
- push 1
- call <f>

3

2

1

retaddr

# gcc Calling Convention (32-bit)

- **Caller is responsible for cleaning arguments off the stack**
  - Note that this allows functions with variable numbers of arguments (varargs)
- **%eax %ecx, %edx may be trashed by callee (must be saved by caller)**
- **%ebp, %ebx, %esi, %edi must be preserved by callee**

# Aside: Buffer Overflow

- Consider a function like:  

```
int f(int arg1, int arg2) {  
    char x[4];  
    ...  
}
```
- Stack frame for function shown at right
- Suppose inside f, we write past the end of array x?
- What gets overwritten? What happens when we execute ret?

Arg2

Arg2

Return address

Saved base pointer

X[3]

X[2]

X[1]

X[0]



# Example from *XV6*: cat.asm

# Example Function Call

```
#include <stdio.h>

int g(int a) {
    return a*2;
}

int f(int a, int b) {
    return g(a) + g(b);
}

int main(void) {
    printf("%d\n", f(1, 2));
    return 0;
}
```

# main

```
int main(void) {  
    printf("%d\n", f(1,2));  
    return 0;  
}
```

08048428 <main>:

```
8048428: 8d 4c 24 04  
804842c: 83 e4 f0  
804842f: ff 71 fc  
8048432: 55  
8048433: 89 e5  
8048435: 51  
8048436: 83 ec 04  
8048439: 6a 02  
804843b: 6a 01  
804843d: e8 c3 ff ff ff  
8048442: 83 c4 08  
8048445: 83 ec 08  
8048448: 50  
8048449: 68 00 85 04 08  
804844e: e8 7d fe ff ff  
8048453: 83 c4 10  
8048456: b8 00 00 00 00  
804845b: 8b 4d fc  
804845e: c9  
804845f: 8d 61 fc  
8048462: c3
```

```
lea    0x4(%esp),%ecx  
and     $0xfffffffff0,%esp  
pushl   -0x4(%ecx)  
push    %ebp  
mov     %esp,%ebp  
push    %ecx  
sub     $0x4,%esp  
push    $0x2  
push    $0x1  
call    8048405 <f>  
add     $0x8,%esp  
sub     $0x8,%esp  
push    %eax  
push    $0x8048500  
call    80482d0 <printf@plt>  
add     $0x10,%esp  
mov     $0x0,%eax  
mov     -0x4(%ebp),%ecx  
leave  
lea     -0x4(%ecx),%esp  
ret
```

**f**

```
#include <stdio.h>

int g(int a) {
    return a*2;
}

int f(int a, int b) {
    return g(a) + g(b);
}

int main(void) {
    printf("%d\n", f(1,2));
    return 0;
}
```

08048405 <f>:

8048405:	55	push	%ebp
8048406:	89 e5	mov	%esp,%ebp
8048408:	53	push	%ebx
8048409:	ff 75 08	pushl	0x8(%ebp)
804840c:	e8 ea ff ff ff	call	80483fb <g>
8048411:	83 c4 04	add	\$0x4,%esp
8048414:	89 c3	mov	%eax,%ebx
8048416:	ff 75 0c	pushl	0xc(%ebp)
8048419:	e8 dd ff ff ff	call	80483fb <g>
804841e:	83 c4 04	add	\$0x4,%esp
8048421:	01 d8	add	%ebx,%eax
8048423:	8b 5d fc	mov	-0x4(%ebp),%ebx
8048426:	c9	leave	
8048427:	c3	ret	

# g

```
#include <stdio.h>

int g(int a) {
    return a*2;
}

int f(int a, int b) {
    return g(a) + g(b);
}

int main(void) {
    printf("%d\n", f(1,2));
    return 0;
}
```

080483fb <g>:

80483fb:	55
80483fc:	89 e5
80483fe:	8b 45 08
8048401:	01 c0
8048403:	5d
8048404:	c3

push	%ebp
mov	%esp,%ebp
mov	0x8(%ebp),%eax
add	%eax,%eax
pop	%ebp
ret	

# Walkthrough

- A great way to understand what's going on is to step through a program in **gdb**
- You'll also use **gdb** to debug the OS, so it's good to become familiar with it
- On OS X the default debugger is **lldb**, which has different syntax for commands (it is possible to install **gdb** instead, though)

# Today: Machine Programming Basics

- **Assembly**
  - More examples
  - Function calls
- **GDB**
- **Booting the OS**
  - BIOS

# Using gdb

- **break <location>**
  - Stop execution at function name or address
  - Reset breakpoints when restarting gdb
- **run <args>**
  - Run program with args <args>
  - Convenient for specifying text file with answers
- **disas <fun>, but not dis**
- **stepi / nexti**
  - Steps / does not step through function calls



# Using gdb

- **info registers**

- Print hex values in every register

- **print (/x or /d) \$eax - Yes, use \$**

- Print hex or decimal contents of %eax

- **x \$register, x 0xaddress**

- Prints what's in the register / at the given address
- By default, prints one word (4 bytes)
- Specify format: /s, /[num][size][format]
  - x/8a 0x15213
  - x/4wd 0xdeadbeef

## GDB

- Use gdbtui
  - Nice display for viewing source/executing commands

```
hello.c
1  #include <stdio.h>
2  #include <unistd.h>
3
4  int main(void)
5  {
6      int i = 1;
7
8      while (i < 60) {
9          i++;
10         sleep(1);
11     }
12
13     return 0;
14 }
15
16

0x8048384 <main>      lea    0x4(%esp),%ecx
0x8048388 <main+4>      and    $0xffffffff0,%esp
0x804838b <main+7>      pushl  -0x4(%ecx)
0x804838e <main+10>     push   %ebp
0x804838f <main+11>     mov    %esp,%ebp
0x8048391 <main+13>     push   %ecx
0x8048392 <main+14>     sub    $0x14,%esp
B+> 0x8048395 <main+17> movl    $0x1,-0x8(%ebp)
0x804839c <main+24>     jmp    0x80483ae <main+42>
0x804839e <main+26>     incl   -0x8(%ebp)
0x80483a1 <main+29>     sub    $0xc,%esp
0x80483a4 <main+32>     push   $0x1
0x80483a6 <main+34>     call   0x80482b8 <sleep@plt>
0x80483ab <main+39>     add    $0x10,%esp
0x80483ae <main+42>     cmpl   $0x3b,-0x8(%ebp)
0x80483b2 <main+46>     jle    0x804839e <main+26>
0x80483b4 <main+48>     mov    $0x0,%eax

child process 9865 In: main                               Line: 6    PC: 0x8048395

Copyright (C) 2008 Free Software Foundation, Inc.
License GPLv3+: GNU GPL version 3 or later <http://gnu.org/licenses/gpl.html>
This is free software: you are free to change and redistribute it.
There is NO WARRANTY, to the extent permitted by law.  Type "show copying"
and "show warranty" for details.
This GDB was configured as "i486-slackware-linux"...
(gdb) b main
Breakpoint 1 at 0x8048395: file hello.c, line 6.
(gdb) r
Starting program: /home/beej/hello

Breakpoint 1, main () at hello.c:6
(gdb) █
```

# DEMO

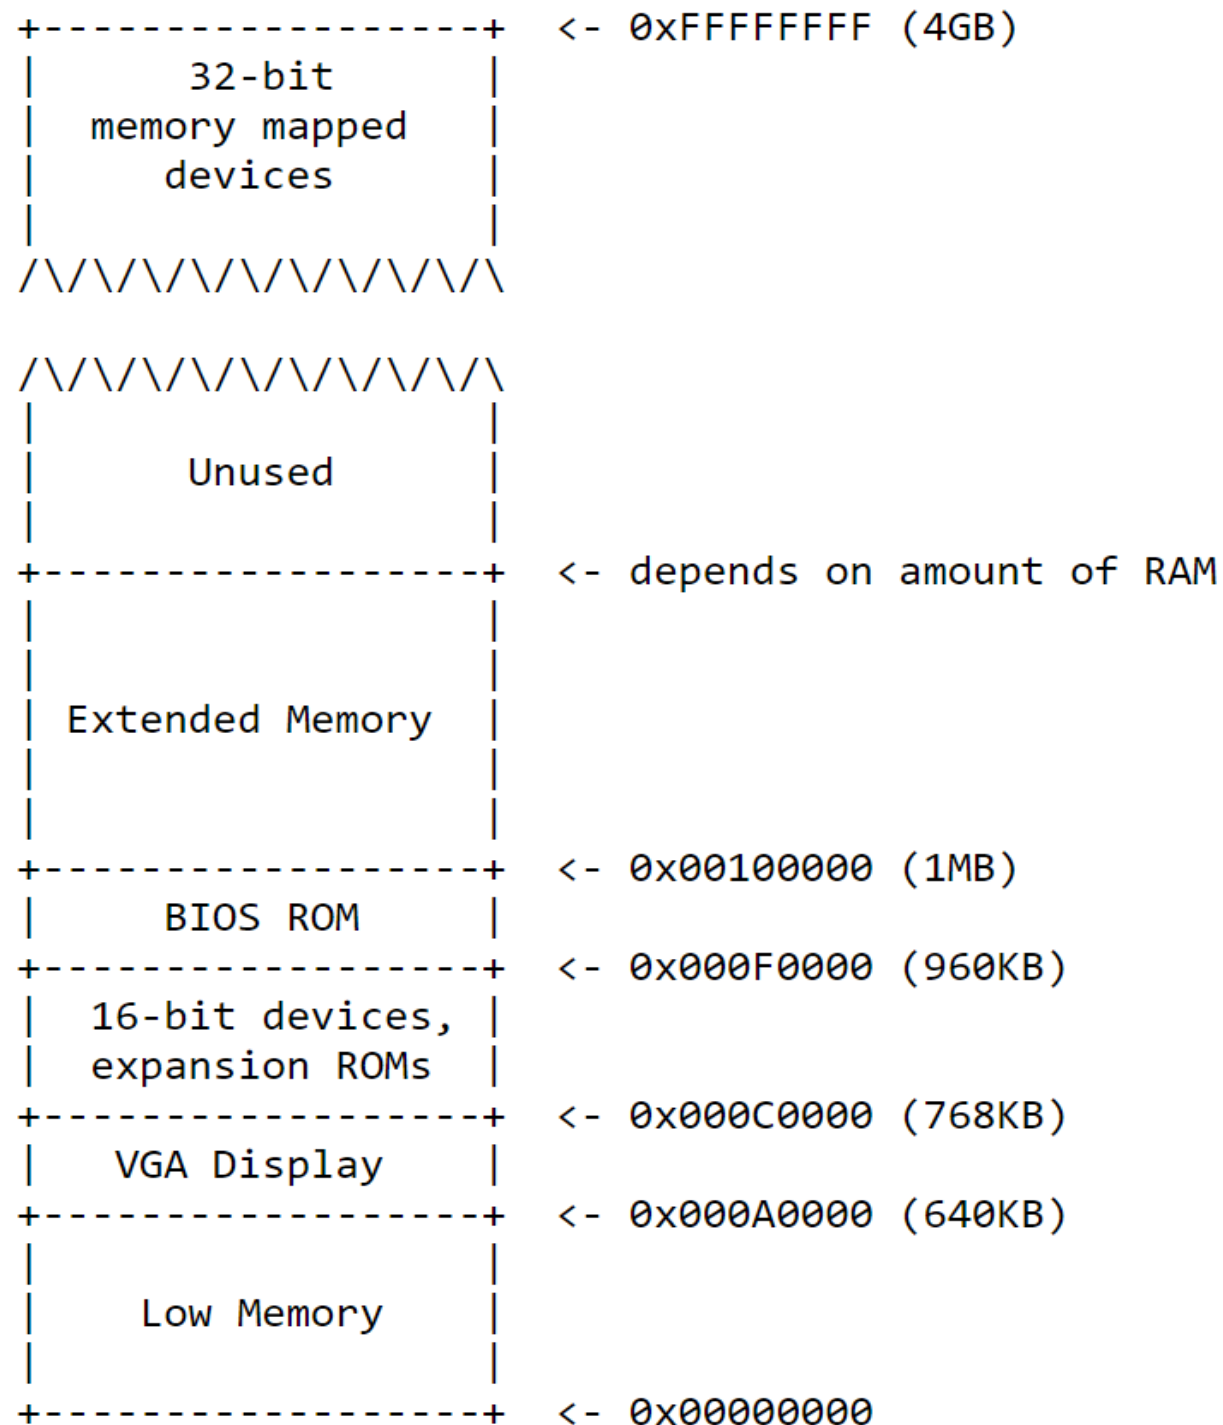
# Today: Machine Programming Basics

- Assembly
  - More examples
  - Function calls
- GDB
- Booting the OS
  - **BIOS**
  - Assembly Bootstrap
  - C Bootstrap

# BIOS

- When an x86 PC boots, it starts executing the **BIOS**.
- **BIOS**
  - stored on the motherboard
  - Prepares the hardware
  - Transfers control to the OS (Bootsector)
- **BootSector:**
  - First 512 bytes of the boot disk.
  - Contains the boot loader

# PC Physical Address Space



# xv6 Bootloader

- Note: x86 processors start out in 16-bit *real mode* (pretends to be an 8088 processor from 1979)
- BIOS loads the bootloader (512 bytes) at 0x7c00
- Loader's job:
  - Put the processor in a more modern operating mode(386 from 8088),
  - Load the xv6 kernel from disk into memory,
  - Transfer control to the kernel.

- The xv6 boot loader comprises two parts:
  - **Assembly Bootstrap(`bootasm.S`):** written in a combination of 16-bit and 32-bit x86 assembly
  - **C Bootstrap(`bootmain.c`);** written in C



# Today: Machine Programming Basics

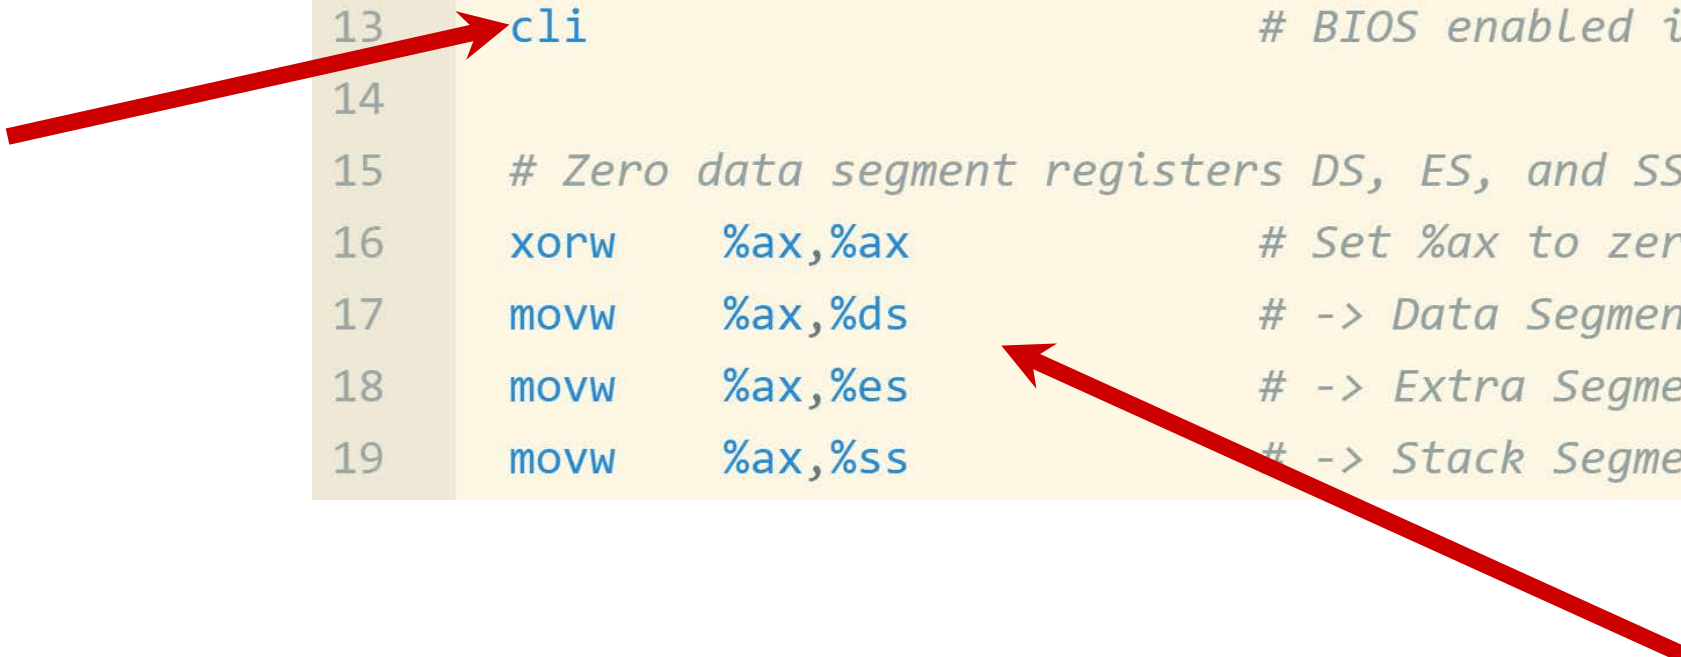
- Assembly
  - More examples
  - Function calls
- GDB
- Booting the OS
  - BIOS
  - **Assembly Bootstrap**
  - C Bootstrap

# Preliminaries

- **During the BIOS's hardware initialization, it enabled hardware interrupts**
  - We don't know what its interrupt handlers look like, and we don't want to deal with interrupts yet
  - Use **cli** instruction to disable interrupts.
- **We also don't know the state of segment registers, so we clear them**

# Preliminaries

```
10  .code16                                # Assemble for 16-bit mode
11  .globl start
12  start:
13  cli                                    # BIOS enabled interrupts; disable
14
15  # Zero data segment registers DS, ES, and SS.
16  xorw    %ax,%ax                        # Set %ax to zero
17  movw    %ax,%ds                        # -> Data Segment
18  movw    %ax,%es                        # -> Extra Segment
19  movw    %ax,%ss                        # -> Stack Segment
```



# The Ugly A20 Hack

- At boot, one of the memory address lines (A20) is forced to 0
- This is for backward compatibility with the original 8088/8086, where addresses were not supported
- So an unused pin on the keyboard controller was attached to the A20 line

# I/O Port for Keyboard Status

# Physical address line A20 is tied to zero so that the first PCs  
# with 2 MB would run software that assumed 1 MB. Undo that.

seta20.1:

inb        \$0x64,%a1                    # Wait for not busy  
testb     \$0x2,%a1  
jnz        seta20.1

movb      \$0xd1,%a1                    # 0xd1 -> port 0x64  
outb      %a1,\$0x64

seta20.2:

inb        \$0x64,%a1                    # Wait for not busy  
testb     \$0x2,%a1  
jnz        seta20.2

movb      \$0xdf,%a1                    # 0xdf -> port 0x60  
outb      %a1,\$0x60

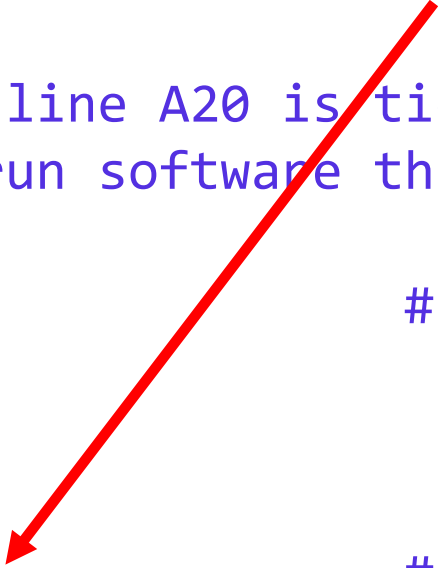
# Write next byte to Controller Output Port

```
# Physical address line A20 is tied to zero so that the first PCs
# with 2 MB would run software that assumed 1 MB. Undo that.
seta20.1:
    inb    $0x64,%a1                # Wait for not busy
    testb  $0x2,%a1
    jnz    seta20.1

    movb   $0xd1,%a1                # 0xd1 -> port 0x64
    outb   %a1,$0x64

seta20.2:
    inb    $0x64,%a1                # Wait for not busy
    testb  $0x2,%a1
    jnz    seta20.2

    movb   $0xdf,%a1                # 0xdf -> port 0x60
    outb   %a1,$0x60
```



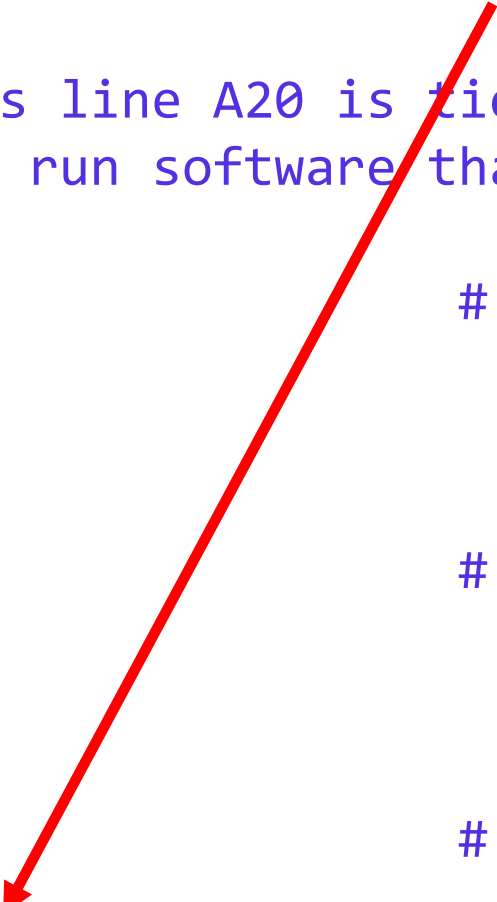
# PS/2 Data Port

```
# Physical address line A20 is tied to zero so that the first PCs
# with 2 MB would run software that assumed 1 MB. Undo that.
seta20.1:
    inb    $0x64,%a1                # Wait for not busy
    testb  $0x2,%a1
    jnz    seta20.1

    movb   $0xd1,%a1                # 0xd1 -> port 0x64
    outb   %a1,$0x64

seta20.2:
    inb    $0x64,%a1                # Wait for not busy
    testb  $0x2,%a1
    jnz    seta20.2

    movb   $0xdf,%a1                # 0xdf -> port 0x60
    outb   %a1,$0x60
```



# A20 Enable

0xdf = 11011111



# Physical address line A20 is tied to zero so that the first PCs  
# with 2 MB would run software that assumed 1 MB. Undo that.

seta20.1:


```
inb    $0x64,%a1    # Wait for not busy
testb  $0x2,%a1
jnz    seta20.1
```

```
movb   $0xd1,%a1    # 0xd1 -> port 0x64
outb   %a1,$0x64
```

seta20.2:

```
inb    $0x64,%a1    # Wait for not busy
testb  $0x2,%a1
jnz    seta20.2
```

```
movb   $0xdf,%a1    # 0xdf -> port 0x60
outb   %a1,$0x60
```





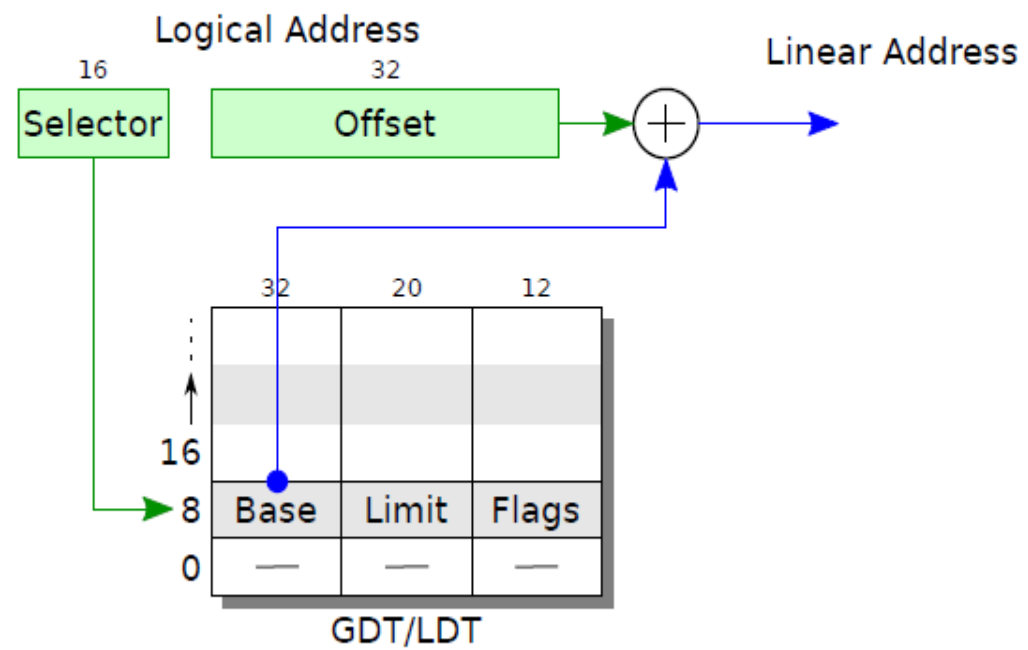
# Transitioning into Protected Mode

- To run 32-bit code, we need to transition into *protected mode*
- A key feature here is that the meaning of the segment registers (%cs, %ds, etc.) changes
- Instead of just holding extra bits for memory addresses, they *select* a particular *segment descriptor*
- *Segment descriptor* gives the base address, size, and permissions of that memory segment

# Segments in xv6

- As with most modern OSes, xv6 makes very little use of segmentation
- The segments it defines just give a 1-1 mapping between logical and physical addresses
- Later, it will set up the *paging hardware* and use *virtual addresses*

# Segments



```
# Switch from real to protected mode.  
# Use a bootstrap GDT that makes  
# virtual addresses map directly to  
# physical addresses so that the  
# effective memory map doesn't change  
# during the transition.
```

```
lgdt    gdtdesc  
movl    %cr0, %eax  
orl     $CR0_PE, %eax  
movl    %eax, %cr0
```

# Protected Mode

- Finally we can jump into *protected mode*. The change happens only once a segment register is modified, so we execute a *long jump* to set %cs
- From here on out, we're executing 32-bit x86 code

```
ljmp    $(SEG_KCODE<<3), $start32
```

```

.code32 # Tell assembler to generate 32-bit code now.
start32:
    # Set up the protected-mode data segment registers
    movw    $(SEG_KDATA<<3), %ax    # Our data segment selector
    movw    %ax, %ds                # -> DS: Data Segment
    movw    %ax, %es                # -> ES: Extra Segment
    movw    %ax, %ss                # -> SS: Stack Segment
    movw    $0, %ax                # Zero segments not ready for use
    movw    %ax, %fs                # -> FS
    movw    %ax, %gs                # -> GS

```

# Onward to C

- Finally, we can set up the stack and jump into C code
- The stack pointer is set to **0x7c00** – the address of the start of the bootloader

```
# Set up the stack pointer and call into C.  
movl    $start, %esp  
call    bootmain
```

# Today: Machine Programming Basics

- Assembly
  - More examples
  - Function calls
- GDB
- Booting the OS
  - BIOS
  - Assembly Bootstrap
  - **C Bootstrap**



# bootmain.c

- Reads in the kernel header from the hard disk
- Parses the kernel headers – it's a standard ELF (Executable and Linkable Format) file that gcc emits
- Reads each section (kernel code, data) into memory at the address specified in the ELF headers
- Finally, calls the *entry point* defined in the ELF file

## ■ Bootmain.c

```
17 void
18 bootmain(void)
19 {
20     struct elfhdr *elf;
21     struct proghdr *ph, *eph;
22     void (*entry)(void);
23     uchar* pa;
24
25     elf = (struct elfhdr*)0x10000; // scratch space
26
27     // Read 1st page off disk
28     readseg((uchar*)elf, 4096, 0);
29
30     // Is this an ELF executable?
31     if(elf->magic != ELF_MAGIC)
32         return; // let bootasm.S handle error
33
34     // Load each program segment (ignores ph flags).
35     ph = (struct proghdr*)((uchar*)elf + elf->phoff);
36     eph = ph + elf->phnum;
37     for(; ph < eph; ph++){
38         pa = (uchar*)ph->paddr;
39         readseg(pa, ph->filesz, ph->off);
40         if(ph->memsz > ph->filesz)
41             stosb(pa + ph->filesz, 0, ph->memsz - ph->filesz);
42     }
43
44     // Call the entry point from the ELF header.
45     // Does not return!
46     entry = (void(*) (void))(elf->entry);
47     entry();
48 }
```

# kernel.ld

- The kernel will eventually be loaded at 0x80100000 (virtual)
- For now though, we have only physical memory, not virtual memory
- Luckily, we can use a *linker script* to specify the physical load address, along with the entry point

```

/* Simple linker script for the JOS kernel.
   See the GNU ld 'info' manual ("info ld") to learn the syntax. */

OUTPUT_FORMAT("elf32-i386", "elf32-i386", "elf32-i386")
OUTPUT_ARCH(i386)
ENTRY(_start)

SECTIONS
{
    /* Link the kernel at this address: "." means the current address */
    /* Must be equal to KERNLINK */
    . = 0x80100000;

    .text : AT(0x100000) {
        *(.text .stub .text.* .gnu.linkonce.t.*)
    }

```

Note that  $0x80100000 = 0x100000 + 0x80000000$

So during early boot we can translate between "virtual" and physical addresses by simple addition & subtraction

# Reading from the Hard Drive

- **You can talk to an IDE hard drive using port I/O (PIO)**
  - This mode is not very fast, but it's simple to implement
  - An alternative is DMA (Direct Memory Access), which allows data transfer to bypass the CPU
  - xv6 only implements PIO mode

```

// Read a single sector at offset into dst.
void
readsect(void *dst, uint offset)
{
    // Issue command.
    waitdisk();
    outb(0x1F2, 1);    // count = 1
    outb(0x1F3, offset);
    outb(0x1F4, offset >> 8);
    outb(0x1F5, offset >> 16);
    outb(0x1F6, (offset >> 24) | 0xE0);
    outb(0x1F7, 0x20); // cmd 0x20 - read sectors

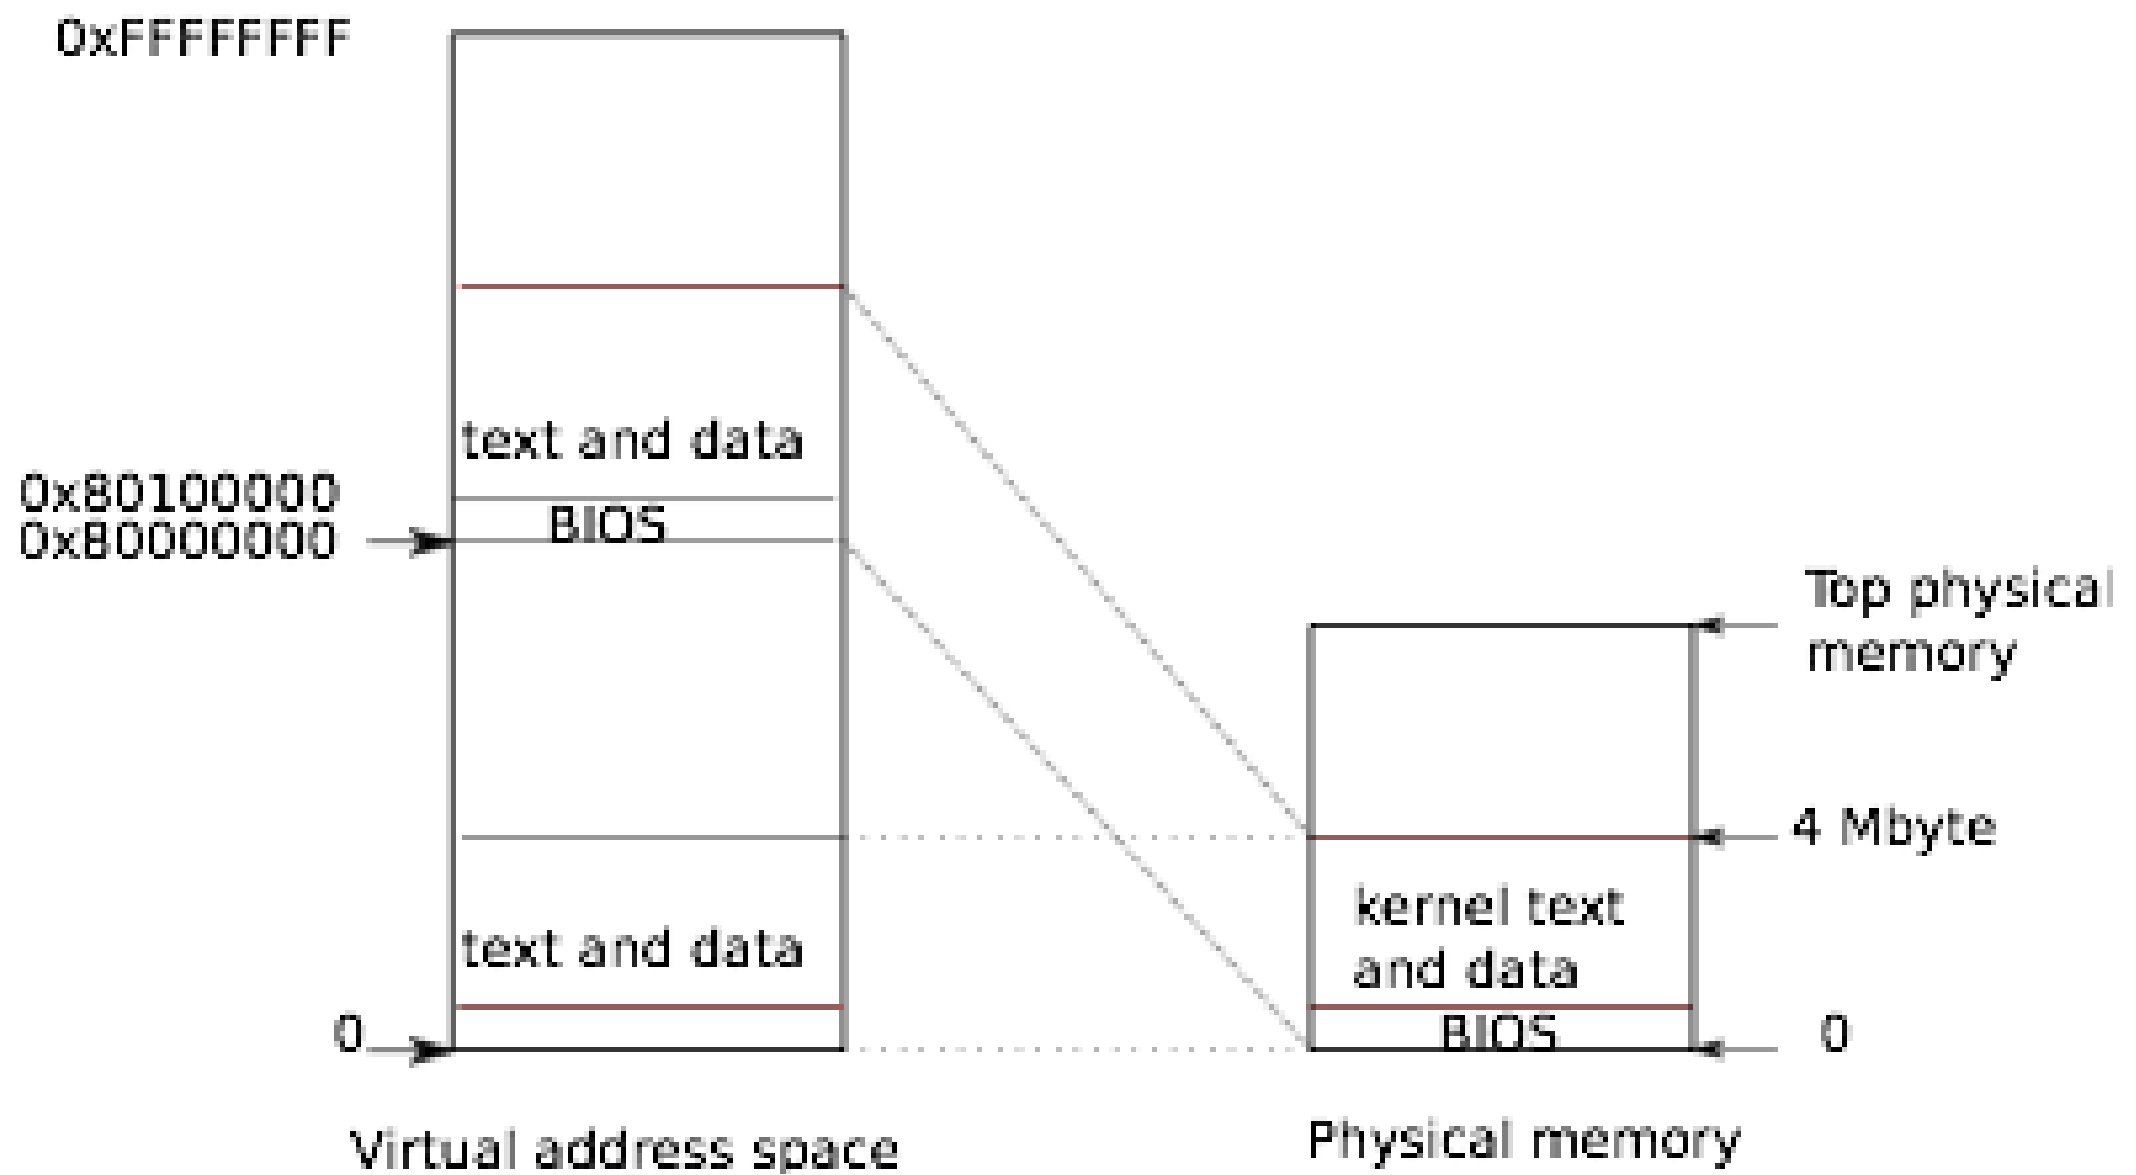
    // Read data.
    waitdisk();
    insl(0x1F0, dst, SECTSIZE/4);
}

```

# Kernel Entry Point

- Defined in entry.S
- This is where we finally turn on paging, but for now we just use a very simple mapping:

**0x80000000 => 0x00000000**





# Entering xv6 on boot processor, with paging off.

`.globl entry`

`entry:`

# Turn on page size extension for 4Mbyte pages

`movl %cr4, %eax`

`orl $(CR4_PSE), %eax`

`movl %eax, %cr4`

# Set page directory

`movl $(V2P_W0(entrypgdir)), %eax`

`movl %eax, %cr3`

# Turn on paging.

`movl %cr0, %eax`

`orl $(CR0_PG|CR0_WP), %eax`

`movl %eax, %cr0`

# Kernel Main

```
// Bootstrap processor starts running C code here.
// Allocate a real stack and switch to it, first
// doing some setup required for memory allocator to work.
int
main(void)
{
    kinit1(end, P2V(4*1024*1024)); // phys page allocator
    kvmalloc(); // kernel page table
    mpinit(); // collect info about this machine
    lapicinit();
    seginit(); // set up segments
    cprintf("\ncpu%d: starting xv6\n\n", cpu->id);
    picinit(); // interrupt controller
    ioapicinit(); // another interrupt controller
    consoleinit(); // I/O devices & their interrupts
    uartinit(); // serial port
    pinit(); // process table
    tvinit(); // trap vectors
    binit(); // buffer cache
    fileinit(); // file table
    ideinit(); // disk
    if(!ismp)
        timerinit(); // uniprocessor timer
    startothers(); // start other processors
    kinit2(P2V(4*1024*1024), P2V(PHYSTOP)); // must come after startothers()
    userinit(); // first user process
    // Finish setting up this processor in mpmain.
    mpmain();
}
```

# Creating the First Process

- This happens in `userinit()` – after we've initialized lots of other hardware [`proc.c:79`]
- The first process will run the code in `initcode.S`
- To do that it needs:
  - A process structure filled out and assigned a slot in the process table (`allocproc()`)
  - An address space (`setupkvm()`)

# initcode.S

```
# exec(init, argv)
.globl start
start:
    pushl $argv
    pushl $init
    pushl $0 // where caller pc would be
    movl $SYS_exec, %eax
    int $T_SYSCALL
[...]
# char init[] = "/init\0";
init:
    .string "/init\0"
```

